
Continuous Fourier-Slice Audits for Cryo-EM Density Uncertainty

Fourier Cryo Splats Project¹

Research manuscript, 29 September 2026.

Under revision after independent review; not submitted.

Abstract

A confidence interval can exclude density errors that lie outside its reconstruction representation. We study this problem for fixed linear density features in cryogenic electron microscopy (cryo-EM), using continuous Fourier-slice adjoints to audit a declared supported L^2 class. The construction specializes classical bias-aware inference. Its operator-specific components are analytic continuous integrals, matrix-free quadrature with explicit integration remainders, and bounds for shared density errors interacting with nonlinear pose uncertainty. All coverage-verification experiments use simulated noise and known simulation limits, with acquisition geometries from three experimental stacks. These tests verify implementations and expose a precision problem: the pose audit can improve on a no-data interval yet have negligible sign-detection power. An integrated-moment refinement reduces the original pose bound’s conservatism, but useful inference at finer structural scales and realistic alignment errors remains unestablished. An experimental sensitivity analysis detects one broad fixed-pose feature after independent noise recalibration, but none at nonzero pose budgets; its nuisance assumptions remain unverified. Reconstructions and background diagnostics provide context, not density-coverage labels. The density radius, support, noise law and pose radii remain assumptions, so the present audit is a conditional research method rather than an experimentally calibrated confidence pipeline.

1. Introduction

Cryo-EM reconstructs molecular density from noisy projections whose orientations, translations and imaging parameters are imperfectly known. Modern neural and Gaussian representations offer expressive reconstruction models, but expressiveness alone does not establish which reconstructed features the data support. We focus on a fixed local density average or contrast. This target differs from a conformational population, the uncertainty of an atomic model, or prediction of a noisy held-out particle image.

An uncertainty calculation usually inherits its reconstruction space. If a Gaussian dictionary spans S , a posterior or confidence interval computed in S need not cover density outside S . Improving optimization inside the dictionary does not remove this discrepancy. Independent half reconstructions can also share the same representational restriction. Our question is how to audit feature uncertainty outside the reconstruction representation. Gaussian pilots motivate the initial exclusion tests; the final continuous solver operates directly on Fourier-slice adjoints and does not require a Gaussian dictionary.

We develop three connected components: a continuous Fourier adjoint-residual audit that separates fitting and inference spaces; bounds for local pose errors, including shared density–pose interactions and an integrated nonlinear remainder; and frozen validation protocols that distinguish fixed-parameter coverage, Bayesian prior-predictive coverage and experimental-image prediction. General bias-aware intervals, norm duality, moment inequalities and primal-dual optimization are established tools. The contribution is their explicit continuous Fourier-slice specialization, its computational implementation, and an empirical audit of when the resulting guarantees remain informative.

2. Related Work and Scientific Scope

Probabilistic reconstruction. Bayesian inference is foundational in cryo-EM (Scheres, 2012). Ullrich et al. (2020) already combine differentiable Fourier-slice reconstruction with variational density uncertainty and discuss the distinction between variance and model bias. Rangan et al. (2024) examine soft modes of the joint volume/alignment likeli-

¹Fourier Cryo Splats research project. Correspondence to: Project repository <<https://github.com/yashpatel5400/fourier-cryo-splats>>.

hood Hessian. Thus neither a probabilistic Fourier representation nor pose–density ambiguity is a new observation. Our candidate emphasizes a finite-sample, bounded-class statement and a check outside the fitting dictionary.

Representations and heterogeneity. CryoDRGN learns a latent-conditioned density family (Zhong et al., 2021a); its latent variation is not automatically calibrated uncertainty in a homogeneous density. Ab initio neural inference (Zhong et al., 2021b) is a different task from fixed-pose backprojection. Gaussian mixtures and Gaussian splatting are also established (Chen & Ludtke, 2021; Chen & Ling, 2026; Qu et al., 2025; Shekarforoush et al., 2025). We use Fourier kernels as a numerical basis, without treating their existence as the contribution. Continuous-state recovery, conformational populations and pose entanglement are substantial separate problems; the living survey records current RECOVAR, DynaMight, cryoDRGN-AI and CAHRA work with source-level distinctions.

Validation and map confidence. Gold-standard splitting and noise substitution address reconstruction overfitting (Scheres & Chen, 2012; Chen et al., 2013); a shell correlation is not a coverage statement for a local feature. LocScale-2.0 provides calibrated confidence for map enhancement relative to its reference target (Bharadwaj et al., 2026). That target must not be conflated with confidence in an unknown density from raw particles. Map significance relative to background noise, directional resolution, atomic-model reliability and particle-derived confidence likewise answer different questions.

Statistical foundations. Optimal recovery links linear inverse estimation to bias–variance tradeoffs (Donoho, 1994). Honest inference over convex function classes and bias-aware Gaussian critical values are classical (Armstrong & Kolesár, 2018). Our theorems below specialize these principles and identify the assumptions that cryo-EM introduces. The survey and its bibliographic ledger are ongoing; a downloaded paper is not counted as a completed critical review.

3. Forward Model and Inference Target

With $F(k) = \int \rho(x) e^{-2\pi i k^\top x} dx$, a particle obeys

$$Y_i(q) = C_i(q) e^{-2\pi i q^\top t_i} F(E_i q) + \epsilon_i(q). \quad (1)$$

E_i embeds a detector plane, C_i is its contrast transfer function (CTF), and frequencies and translations use reciprocal units. We realify one representative of each conjugate Fourier pair and whiten the noise. The initial theory conditions on a fixed design and an independent pilot ρ_0 , giving

$$y - \mathcal{A}\rho_0 = \mathcal{A}\delta + Ju + r + \epsilon, \quad \epsilon \sim N(0, I). \quad (2)$$

Here $\delta = \rho - \rho_0$, J is block diagonal over particles, u describes bounded nuisance errors, and r is a bounded remainder. We assume

$$\|\delta\|_H \leq B, \quad \|u_i\| \leq \eta_i, \quad \|r_i\| \leq \gamma_i. \quad (3)$$

H is a declared Gaussian subspace, a supported voxel space, or continuous L^2 on a stated cube. These are different classes. A physical density-energy norm includes voxel volume; experiments using an ordinary voxel Euclidean norm explicitly normalize that norm. B , η_i and γ_i are assumptions, not quantities automatically learned from residual fit.

The estimand is $t = \langle \ell, \rho \rangle$, such as a Gaussian-weighted spatial average or the difference of two such averages. The smoothing scale and coordinate frame are specified in physical units. For weights fixed without the inference noise, use

$$\hat{t} = \langle \ell, \rho_0 \rangle + w^\top (y - \mathcal{A}\rho_0). \quad (4)$$

The pilot, dictionary, whitening, target and weights may depend on independent development data. Poses estimated from the inference images violate literal fixed-design independence; supplied consensus poses therefore support only a conditional experimental analysis here.

4. Bias-Aware Confidence and Ambient Auditing

4.1. A conditional certificate

Define $s = \|w\|$ and

$$b_H(w) = B \|\ell - \mathcal{A}^* w\|_H + \sum_i \eta_i \|J_i^\top w_i\| + \sum_i \gamma_i \|w_i\|. \quad (5)$$

Let $q(s, b, \alpha)$ solve

$$\Phi((q - b)/s) - \Phi((-q - b)/s) = 1 - \alpha \quad (6)$$

when $s > 0$, and set $q(0, b, \alpha) = b$.

Theorem 1 (Conditional coverage). *Under Equations (2)–(3), the interval $\hat{t} \pm q(s, b_H(w), \alpha)$ covers t with probability at least $1 - \alpha$, uniformly over the stated class, conditional on the independent design choices.*

Proof. The deterministic estimation error is $\langle \mathcal{A}^* w - \ell, \delta \rangle + \sum_i w_i^\top J_i u_i + \sum_i w_i^\top r_i$. Cauchy–Schwarz bounds its absolute value by b_H . The remaining error is Gaussian with standard deviation s . Coverage of a symmetric interval decreases with the absolute Gaussian mean, so the minimum occurs at mean $\pm b_H$ and equals Equation (6). $\square$

The result is a direct bias-aware inference construction. It is useful only if its assumptions are defensible and its width improves on the no-data interval $\langle \ell, \rho_0 \rangle \pm B \|\ell\|$.

4.2. Why auditing the dictionary matters

Let $S : \mathbb{R}^p \rightarrow H$ have orthonormal columns. A restricted calculation replaces the density term by $B \|S^*(\ell - \mathcal{A}^*w)\|$. Write $h = \ell - \mathcal{A}^*w$. Then

$$\|h\|_H^2 = \|S^*h\|^2 + \|(I - SS^*)h\|_H^2. \quad (7)$$

The second term measures uncertainty discarded by the dictionary. It can be large even when the restricted residual is zero. Evaluating the full adjoint residual repairs this omission for H without estimating a representation-error direction from noisy data. It does not cover density outside the declared ambient grid or support.

Centering in Equation (4) also uses the ambient operator. If a surrogate centers the images instead, the discrepancy $w^\top(\mathcal{A} - \bar{\mathcal{A}})\rho_0$ requires an additional bound. We do not silently substitute a surrogate forward model into the coverage theorem.

4.3. Optimization and verifiable stopping

We optimize the conservative convex objective

$$\mathcal{F}_H(w) = z_{1-\alpha/2} \|w\| + b_H(w), \quad (8)$$

then report the sharper width from Equation (6). Its dual is

$$\begin{aligned} \max_{v,a,d} \quad & \langle \ell, v \rangle \\ \text{s.t.} \quad & \|v\| \leq B, \quad \|a_i\| \leq \eta_i, \quad \|d_i\| \leq \gamma_i, \\ & \|\mathcal{A}v - Ja - d\| \leq z_{1-\alpha/2}. \end{aligned} \quad (9)$$

Representing each norm by a Euclidean support function proves this dual; continuity and coercivity yield strong duality. A feasible dual iterate supplies a lower bound before convergence. The gap concerns Equation (8), not optimality of the folded-normal width.

For a reduced density space S , a feasible dual vector lifts to Sv in H . Thus a restricted dual lower bound and an ambient primal upper bound bracket the same full-space optimum. We enrich S with the component of $\ell - \mathcal{A}^*w$ orthogonal to it, using only design and target information. Restricted exact optima increase under enrichment; individual ambient primal values need not decrease, so we keep the best feasible iterate. No convergence-rate claim is made for this greedy construction.

Quadratic majorization gives coefficient-space conjugate-gradient solves with particle-block inverses. A classical Chambolle–Pock implementation is an alternative (Chambolle & Pock, 2011). For fixed poses, a full-space matrix-free comparator solves $(\mathcal{A}^*\mathcal{A} + \lambda I)v = \ell$ within majorization. It can outperform enrichment and is included to prevent an unsupported Gaussian speed claim.

5. Continuous Density and Bounded Pose Errors

5.1. A Fourier audit independent of the fitting dictionary

For fixed poses, the audit also extends to $H = L^2([-1/2, 1/2]^3)$ without a voxel or density bandwidth assumption. This cube is a broader support class than the earlier spherical mask. Its Fourier kernel is $K(v) = \prod_a \text{sinc}(v_a)$. If $c_j = (w_{Rj} + iw_{Ij})C_j/\sigma$ and $f_w(x) = \text{Re} \sum_j c_j e^{2\pi i k_j^\top x}$, then

$$\|f_w\|_H^2 = \frac{1}{2} \text{Re} \sum_{j,l} \{c_j \bar{c}_l K(k_j - k_l) + c_j c_l K(k_j + k_l)\}. \quad (10)$$

Gaussian target norms and cross terms have analytic truncated-normal integrals, giving $\|\ell - \mathcal{A}^*w\|_H$ directly. The continuous objective has a feasible dual vector $v = th$, where $h = \ell - \mathcal{A}^*w$ and $t = \min(B/\|h\|, z/\|\mathcal{A}h\|)$. Its lower bound is $t(\|\ell\|^2 - w^\top \mathcal{A} \ell)$.

A matrix-free variant integrates only the finite Fourier adjoint field with Gauss–Legendre quadrature. An explicit quadrature remainder protects both primal and dual bounds; it does not discretize the unknown density class. Standard pivoted Cholesky preconditioning accelerates the solves (Harbrecht et al., 2012). The formulas and error derivation are in Appendix E. A separate continuous nonlinear-pose extension bounds the norm of the first and second polynomial Fourier fields by padded quadrature Grams. Each particle has a normalized five-dimensional pose vector u_i in the joint unit ball: three rotation coordinates with radius a and two detector shifts with radius s_t . This excludes simultaneous maximal rotation and translation. For fixed w , first and second adjoint derivatives form five linear and fifteen symmetric quadratic Fourier fields per particle. Their combination at allowed u_i is bounded by a shared-field norm L_{12} , using spectral and triangle bounds with integration padding. The unknown density is shared across particles; treating it separately in each image would lose that structure. Finite Gaussian/voxel derivations and alternatives appear in Appendix A. With $P = \|\rho_0\|$ and their shared-field bound L_{12} , its bias bound is

$$b \leq B \|h\| + (B + P) \left(L_{12} + \frac{1}{6} \sum_i K_{3i} \|w_i\| \right). \quad (11)$$

Here K_{3i} uniformly bounds the continuous third-derivative operator. Appendix E.1 derives the constants and integration error. The frozen studies post-audit fixed-pose weights. The post-review implementation also optimizes a continuous pose-aware objective and tightens the subsequent audit using residual/pose cross terms and exact known-pilot pairings (Appendix F); it does not claim nonlinear optimality.

An integrated-moment refinement replaces the spatial-maximum cubic remainder by

$$r_{\text{mom}}(w) = \frac{B+P}{6} \sum_{i,q} |c_{iq}| (L_6^3 + 3L_4H_4 + T_2). \quad (12)$$

Here $L_p = 2\pi(a\|k\|c_p + s_t\|q\|)$, $H_p = 2\pi a^2\|k\|c_p$, $T_2 = 2\pi a^3\|k\|c_2$, and $(c_2, c_4, c_6) = (12^{-1/2}, (3/12^2)^{1/4}, (15/12^3)^{1/6})$. These constants bound the cube-integrated moments of linear spatial coordinates. Hölder’s inequality controls the exact third phase derivative; retaining actual frequency weights avoids a Hilbert–Schmidt relaxation (Appendix E.2). No distribution on the unknown density is introduced.

Operationally, we fix the pilot, target, noise model and density/pose classes independently of inference noise; optimize fixed-pose weights; audit their continuous residual and pose derivatives; then use the smallest of the valid remainder bounds. If the resulting interval is wider than the no-data alternative, we switch both center and width to that alternative using design information alone. Only then are inference images used in Equation (4). The moment refinement is evaluated in a separate frozen follow-up, preserving the original study.

6. Experiments

Results below distinguish exploratory development from three completed protocols frozen before their respective outcomes: original continuous uncertainty validation, an integrated-moment follow-up, and additional-exposure prediction. Experimental development pools had already been inspected; only the specified frozen follow-ups constitute prospective tests.

6.1. Data, references and provenance

We retain 8,192 extracted experimental images from each of EMPIAR-10028, 10049 and 10076 (EMPIAR, 2015a;b; 2016). Source exposure groups are reconstructed from deposited metadata. EMPIAR-10076 requires the original Frealign FILM field because its STAR micrograph numbers are sequential particle identifiers. Hash-based whole-group partitions provide pilot, tuning, inference-half and test labels within this development pool. Consensus poses and CTFs remain conditional inputs. EMPIAR-10076 contains heterogeneous assembly intermediates, so a single-density fit is a working consensus model, not experimentally established homogeneity.

Deposited maps EMD-2660, EMD-6487 and EMD-8434 define independent voxel signal generators after explicit low-pass resampling. They are simulation truth, not proof of experimental accuracy. Nonuniform Fourier transforms project these maps independently of the Gaussian represen-

tation. Native and downloaded-file hashes, source indices, signs, field sizes and software versions are recorded.

The physical fields are 482.4, 236.16 and 419.2 Å, respectively. Gaussian feature standard deviations of 0.03/0.07 times the field therefore span 7.1–33.8 Å. The radius-five inference band reaches wavelengths 47.2–96.5 Å; the radius-twelve extension reaches 19.7–40.2 Å. These uncertainty experiments probe coarse structural features, not near-atomic reconstruction. The separate image-prediction comparison uses a wider common band on the 64×64 working grid.

6.2. Matched controls and exploratory finite-space studies

The numerical controls use identical density functionals and declared units: full/diagonal Gaussian posterior covariance, regularized sampling covariance, joint Gaussian pose covariance, source-group bootstrap, fixed-pose bias-aware intervals and the no-data estimator. Posterior methods pass the matched prior-predictive identity; fixed-parameter boundary coverage is a different criterion. These are mathematical controls, not claims to reproduce every external method’s complete implementation.

The exploratory studies in Appendix B motivate the continuous extension. A dictionary can exclude measurable bias even when an interval is calibrated in that dictionary. Quadratic pose terms and shared-density spectral bounds improve a first-order relaxation. Weakly regularized full-voxel bootstrap coverage is 0.938/0.956/0.958 with smaller widths than the robust audit; stronger regularization can reduce coverage to 0.339. Thus the audit does not uniformly dominate empirical interval procedures in width. Grid agreement alone does not prove a continuous-density guarantee.

6.3. Frozen continuous implementation and precision checks

The first protocol was frozen at commit 76dbd53 before its outcomes. Four additional-cohort geometry subsets per stack, two feature scales, two targets and two pose budgets give 96 audit settings and 2,304 signal/scenario records. Each subset contains 128 particles; its prescribed density radius is $B = 2$, and the simulated Gaussian noise scale is fixed from a pilot SNR convention. Four unit-norm cell generators include a deposited map, ellipsoid unions, modulated shells and signed multiscale fields. Exact continuous-density boundaries exercise directions outside the cell representation. The protocol includes nominal, random-boundary and coherent-boundary poses, plus separately reported four-fold pose exceedance and doubled noise scale. Acquisition subsets may share exposure groups; they are not independent biological replicates.

All 48 unique fixed-pose fits meet the 0.5% sum-objective

gap. Across all 1,536 in-class signal/scenario records, minimum analytic coverage, used here as an implementation check, is 0.99377 for the continuous pose audit, 0.93922 for the fixed-pose audit and zero for the 24-cell projection audit. The cell-defined generators are covered conservatively by both audited fixed-pose procedures; the severe cell-audit failures occur at explicitly continuous boundaries. These are worst-case implementation checks, not typical-map failure or experimental calibration estimates. Scalar Gaussian draws provide only a numerical cross-check and are retained in the historical records.

Figure 1 shows all selected designs. Broad-feature relative widths are 0.109–0.175 at 0.1° and 0.271–0.457 at 0.5° . Fine-feature widths rise to 0.759–0.779 and 0.885–0.998, respectively, nearly reaching the no-data alternative. None of the 96 original-audit settings switches to no data. The pose audit’s minimum calculated coverage across the 768 model-violation records is 0.99948. Its conservatism prevents those selected severities from exposing a failure; this supplies no robustness or calibration claim. Sign power, marginal Monte Carlo intervals, all methods and every failure remain in the released records.

6.4. Frozen integrated-moment follow-up

The follow-up was frozen at 8d66785 after the refinement’s development, with fresh signal/noise seeds and particle subsets disjoint from the original study. Its 48 settings use broad targets and $1^\circ/2^\circ$ budgets. All 24 unique fits converge within the 0.5% gap. Across 768 in-class records, the refined audit’s analytic coverage rounds to one in float64, whereas the fixed-pose audit reaches 0.67114 and the cell-projected audit 0.12596. The refined intervals are also too conservative for the 384 selected model-violation records to expose failure.

The spatial-maximum bound uses no data in all 48 settings; integrated moments reduce this to 16 (Table 1). All 1° widths are 0.304–0.501 of no data. At 2° , only EMPIAR-10076 remains narrower, at 0.866–0.987. However, correct-sign power for the in-class cell generators is at most 2.3×10^{-5} ; constructed continuous boundaries attain at most 0.5184. Thus the gain in interval width is real but does not demonstrate useful detection in these map-like signals. Figure 7 and released records retain all settings.

6.5. Continuous audit and assumption failures

The continuous fixed-pose comparison uses the same three acquisition geometries, two targets and two averaging scales. Pilots and independent map generators are orthonormal constant-cell functions on grids 24 and 64, respectively. Their exact Fourier transforms include cell sinc factors and a half-cell phase; target integrals use normal-CDF differences. Both have unit continuous L^2 norm, so the same

Table 1. Frozen continuous-density and pose validation: relative half-width ranges over both targets and all four geometry replicates per stack. All 48 settings and all in-class signals are retained. The density, noise and pose limits are supplied assumptions.

EMPIAR	Scale 0.07	
	1°	2°
10028	0.39–0.50	1.00–1.00
10049	0.38–0.47	1.00–1.00
10076	0.30–0.36	0.87–0.99

prescribed $B = 2$ covers their difference. The inference class itself contains all signed L^2 functions on the cube; no positivity or atomic-structure restriction is imposed.

All 12 dense continuous solves meet the 0.5% sum-objective gap, and independently evaluated matrix-free results agree to numerical tolerance. Figure 11 distinguishes the continuous bound from projections onto finite cell spaces. At a constructed continuous bias boundary, small relative width differences can matter: a 64-cell interval at the EMPIAR-10028 broad central estimator has coverage approximately 0.923 despite a width difference of only 1.7%. The independent map’s coverage remains one; this adversary is not a typical-signal claim. Refitting weights for the larger cube class is essential: transferring sphere-only weights can be much less informative.

A classical Gaussian two-point lower bound supplies an additional design-only efficiency check (Corollary 1). Across all 18 low- and higher-band fixed-pose targets, its necessary half-length is 0.840–0.882 of the reported interval half-width. Thus the fixed-pose procedures are within a factor 1.20 of the minimum possible deterministic interval length under the stated continuous class, subject to the reported numerical accuracy. This compares against arbitrary interval centers, but not variable-length intervals. The factor is consistent with the classical testing/duality connection below; it is not a new minimax principle. It does not imply that the additional pose bound is comparably sharp.

The first 128-particle timing comparison uses about 419 MB for dense Gram arrays, versus 86 MB for quadrature plus a 1,024-column preconditioner. Setup plus solve/audit takes approximately 14 versus 6 seconds under concurrent Mac load; these stored-array counts are not peak resident memory. An initial weak preconditioner was substantially slower and is retained. The matrix-free implementation also completed both broad targets with 1,024 particles for all three geometries. A separate radius-12 continuous comparison completes both fine targets on all three geometries: relative widths range from 0.120 to 0.238, with all six sum-objective gaps below 0.5%. Individual solve/audit times range from 483 to 1,223 seconds. Despite improvement over no data, reference-generator sign power is below 1.1×10^{-9} in every fine-target case. These results do not establish biologically




Figure 1. Completed frozen original-audit study: all 96 settings. Each geometry contributes both feature targets at each scale and pose budget. Width is relative to the same density class’s no-data interval. The right panel reports minimum analytic coverage by generator/scenario group; “L2” denotes densities at the declared ball boundary, chosen in the continuous adjoint-residual direction to maximize bias at a fixed pose. Values are rounded for display. Fine-feature widths near one and conservative cell-generator coverage limit a usefulness interpretation.

useful fine-feature detection. Projecting the same residual onto 64-cell densities yields coverage 0.791–0.894 at the constructed continuous boundary, despite widths only 3.7–6.5% smaller.

The joint continuous-pose post-audit is complete for both broad targets, all three geometries and rotation budgets $0.1^\circ/0.5^\circ$, with a joint shift budget of 0.01/24 field units (0.201, 0.098 and 0.175 Å, respectively). These are very small translation sensitivities, not empirically calibrated alignment errors. Its twelve relative half-widths range from 0.110 to 0.444 (Figure 12). Exact nonlinear constant-cell projections give reference coverage at least 0.9999999998 in these tests, revealing conservatism rather than empirical calibration. At one nonzero coherent pose per case, analytic sinc integrals eliminate the entire continuous density ball; its feasible bias is 0.24–0.86 of the uniform upper bound. These are lower bounds on worst bias, not globally optimized adversaries. Each audit takes about 38–40 seconds and stores approximately 725 MB of field/Gram arrays under concurrent Mac load.

Six explicit assumption-stress runs each retain 273 records (Figure 13). Support-only violations preserve the full density-energy radius, isolating the damage caused by an incorrect mask. View-coupled two-state mixtures preserve an equal-population mean but violate homogeneity. Both can cause severe undercoverage, as can correlated noise and understated noise scale. The tested low-bandwidth defocus offsets and elliptical Student- t controls do not cause undercoverage for the shared-density pose bounds. Every severity setting and unfavorable result is retained.

A wider development sweep exposes the original continuous

bound’s limits: every broad target falls back to no data at 1° , 2° and 5° . The integrated-moment refinement reduces cubic remainder bounds by 6.79–14.74 times across all 30 small- and large-angle settings, without changing weights. No-data fallbacks decrease from 18 to 10; all six 1° intervals have relative widths 0.300–0.532 (Figure 8). This is post-hoc development, not the frozen study’s outcome. The remaining gap to feasible nonlinear biases and the large-angle failures are retained.

A stronger continuous stress search uses three starts for each bias sign, both targets, all three geometries and $0.5^\circ/1^\circ/2^\circ$ budgets. All 108 final candidates are recomputed with analytic continuous-density integrals and exact nonlinear poses in float64 after an accelerated quadrature-surrogate search. The strongest feasible bias in each setting is 1.06–2.54 times its coherent-pose control, but only 0.16–0.57 of the integrated upper bound (Figure 9). None exceeds that bound. These local searches expose remaining slack; they neither prove a global maximum nor supply validated floating-point lower bounds.

6.6. Baselines and numerical checks

Implemented matched mathematical baselines include full Gaussian posterior covariance, diagonal variational covariance, regularized sampling covariance, joint Gaussian pose covariance, identifiable unregularized estimation, fixed-pose bias-aware intervals, coarse and structured nuisance bounds, and the no-data interval. These are not executions of every external paper’s original code. Full Gaussian posteriors pass their exactly matched linear prior-predictive identity at 0.95. Narrow priors can nevertheless fail bounded fixed-signal

tests; wider priors and unregularized estimation can produce very wide intervals.

Independent conic optimization checks the majorization, primal-dual and ambient solvers on small instances. Derivative finite differences, nonlinear remainder tests, direct Fourier sums, adjoint identities and independent Gaussian-integral checks validate implementation components. The passing test suite includes a three-dimensional Fourier-transform check for externally generated maps. Tight FINUFFT tolerances are numerical checks, not interval-arithmetic guarantees on roundoff.

Post-review development. A matrix-free shared-field operator and a separately budgeted randomized spectral upper audit now permit larger pose calculations. A new convex pose-aware objective passes a small independent conic check; its large solves are still under development. An experimental covariance-envelope attempt on all three stacks initially detects no feature. An independent calibration of the same estimators’ directional noise reduces standard-deviation bounds to 0.132–0.285 of their original values. It detects one broad fixed-pose feature, but none at nonzero pose budgets. Declared CTF errors make the higher-band triangle audit largely vacuous. Appendix F preserves these methods, completed negative outcomes and remaining assumptions.

7. Limitations

The current evidence does not establish an end-to-end uncertainty method with estimated poses, noise law and density class. Local pose radii and the total density bound are assumptions. Selecting them from the same inference residuals would require new analysis; a narrow residual does not rule out an almost invisible density error. Learned support can exclude true structure. Gauge and hand ambiguity must be fixed or included in the target class.

Our three stacks supply distinct experimental geometries, not ground-truth density coverage labels. The uncertainty simulations use coarse features and prescribed noise; the total density ball includes signed functions and is deliberately broader than molecular structure. Its fixed-pose near-optimality result therefore measures performance in that class, not optimality under positivity or atomic constraints. Pose corrections remain loose, and low sign power limits structural use even when widths improve over no data. Source-group splits do not undo full-data consensus-pose dependence. The shared-density assumption is also unverified: EMPIAR-10076 contains ribosomal assembly intermediates (EMPIAR, 2016). Arbitrary calibration means do not extend the inference theorem to view-dependent structural mixtures. Full external probabilistic reconstruction pipelines, joint alignment resampling and converged ab initio comparisons

remain unexecuted; mathematical covariance controls are narrower comparisons. The supplied class assumptions and empirical background correlations prevent an end-to-end calibration claim. Simultaneous maps require multiplicity control, and the targeted literature survey is not exhaustive across all cryo-EM applications.

Reproducibility and Impact

The active research branch is `uncertainty-study` at <https://github.com/yashpatel5400/fourier-cryo-splats>. It preserves failed development approaches and distinguishes them from completed outputs. The earlier reconstruction release v0.1.0 remains available with its conditional FSC results; it is not the uncertainty study’s accuracy claim. No new conference submission or acceptance is implied by this format. Background PDFs are retained locally with source hashes rather than redistributed. Prematurely calibrated-looking maps can mislead structural interpretation; explicit assumptions and negative controls are central outputs of this work.

References

- Armstrong, T. B. and Kolesár, M. Optimal Inference in a Class of Regression Models. *Econometrica*, 86(2):655–683, 2018. doi: 10.3982/ECTA14434.
- Battle, P., Patil, P., Stanley, M., Ruiz Lupon, J., Owahdi, H., and Kuusela, M. Simultaneous Frequentist Calibration of Confidence Regions for Multiple Functionals in Constrained Inverse Problems. *arXiv:2510.11708v1*, 2025. URL <https://arxiv.org/abs/2510.11708v1>.
- Beckers, M., Jakobi, A. J., and Sachse, C. Thresholding of cryo-EM density maps by false discovery rate control. *IUCrJ*, 2019. doi: 10.1107/S2052252518014434.
- Bharadwaj, A., de Bruin, R., and Jakobi, A. J. Confidence-guided cryo-EM map optimisation with LocScale-2.0. *Nature Communications*, 2026. doi: 10.1038/s41467-026-75327-8.
- Cai, T. T. and Low, M. G. An adaptation theory for nonparametric confidence intervals. *The Annals of Statistics*, 32(5):1805–1840, 2004. doi: 10.1214/009053604000000049. URL <https://arxiv.org/abs/math/0503662>.
- Chambolle, A. and Pock, T. A First-Order Primal-Dual Algorithm for Convex Problems with Applications to Imaging. *Journal of Mathematical Imaging and Vision*, 40:120–145, 2011. doi: 10.1007/s10851-010-0251-1.
- Chen, M. and Ludtke, S. J. Deep learning-based mixed-dimensional Gaussian mixture model for characterizing

- variability in cryo-EM. *Nature Methods*, 2021. doi: 10.1038/s41592-021-01220-5.
- Chen, S. and Ling, H. CryoSplat: Gaussian Splatting for Cryo-EM Homogeneous Reconstruction. In *ICLR*, 2026. URL <https://arxiv.org/abs/2508.04929>.
- Chen, S., McMullan, G., Faruqi, A. R., Murshudov, G. N., Short, J. M., Scheres, S. H. W., and Henderson, R. High-resolution noise substitution to measure overfitting and validate resolution in 3D structure determination by single particle electron cryomicroscopy. *Ultramicroscopy*, 135: 24–35, 2013. doi: 10.1016/j.ultramic.2013.06.004.
- Chung, S.-C. Cryo-forum: A framework for orientation recovery with uncertainty measure with the application in cryo-EM image analysis. *arXiv:2307.09847v1*, 2023. URL <https://arxiv.org/abs/2307.09847v1>.
- Dai, M., Dong, Z., Xu, K., and Zhang, Q. C. CryoRes: Local Resolution Estimation of Cryo-EM Density Maps by Deep Learning. *Journal of Molecular Biology*, 2023. doi: 10.1016/j.jmb.2023.168059.
- Donoho, D. L. Statistical Estimation and Optimal Recovery. *The Annals of Statistics*, 22(1):238–270, 1994. doi: 10.1214/aos/1176325367.
- EMPIAR. EMPIAR-10028: Cryo-EM structure of the Plasmodium falciparum 80S ribosome, 2015a. URL <https://www.ebi.ac.uk/empiar/EMPIAR-10028/>.
- EMPIAR. EMPIAR-10049: Cryo-EM structures of synaptic RAG1-RAG2 complex, 2015b. URL <https://www.ebi.ac.uk/empiar/EMPIAR-10049/>.
- EMPIAR. EMPIAR-10076: L17-depleted 50S ribosomal intermediates, 2016. URL <https://www.ebi.ac.uk/empiar/EMPIAR-10076/>.
- Evans, L., Dingeldein, L., Covino, R., Gilles, M. A., Thiede, E., and Cossio, P. Counting particles in cryo-electron microscopy may result in incorrect population estimates. *Communications Biology*, 2026. doi: 10.1038/s42003-026-09859-6.
- Gilles, M. A. and Singer, A. Cryo-EM heterogeneity analysis using regularized covariance estimation and kernel regression. *Proceedings of the National Academy of Sciences of the United States of America*, 2025. doi: 10.1073/pnas.2419140122.
- Harbrecht, H., Peters, M., and Schneider, R. On the low-rank approximation by the pivoted Cholesky decomposition. *Applied Numerical Mathematics*, 62(4):428–440, 2012. doi: 10.1016/j.apnum.2011.10.001.
- Hoff, S. E., Thomasen, F. E., Lindorff-Larsen, K., and Bonomi, M. Accurate model and ensemble refinement using cryo-electron microscopy maps and bayesian inference. *PLOS Computational Biology*, 20(7):e1012180, 2024. doi: 10.1371/journal.pcbi.1012180.
- Janson, A. and Andén, J. Moment-based posterior sampling for multi-reference alignment. *arXiv preprint arXiv:2510.12651*, 2025. URL <https://arxiv.org/abs/2510.12651v1>.
- Jeon, M., Raghu, R., Astore, M., Woollard, G., Feathers, R., Kaz, A., Hanson, S. M., Cossio, P., and Zhong, E. D. CryoBench: Diverse and challenging datasets for the heterogeneity problem in cryo-EM. *arXiv:2408.05526v2*, 2025. URL <https://arxiv.org/abs/2408.05526v2>.
- Kimanius, D., Jamali, K., Wilkinson, M. E., Lövestam, S., Velazhahan, V., Nakane, T., and Scheres, S. H. W. Data-driven regularization lowers the size barrier of cryo-EM structure determination. *Nature Methods*, 2024. doi: 10.1038/s41592-024-02304-8.
- Kuczynski, J. and Wozniakowski, H. Estimating the largest eigenvalue by the power and lanczos algorithms with a random start. *SIAM Journal on Matrix Analysis and Applications*, 13(4):1094–1122, 1992. doi: 10.1137/0613066.
- Lai, T. L., Wang, S.-H., Chung, S.-C., Chang, W.-h., and Tu, I.-P. Uncertainty Quantification in Dynamic Image Reconstruction with Applications to Cryo-EM. *Statistica Sinica*, 33:1771–1788, 2023. doi: 10.5705/ss.202021.0419.
- Levy, A., Raghu, R., Feathers, J. R., Grzadkowski, M., Poitevin, F., Johnston, J. D., Vallese, F., Clarke, O. B., Wetzstein, G., and Zhong, E. D. CryoDRGN-AI: neural ab initio reconstruction of challenging cryo-EM and cryo-ET datasets. *Nature methods*, 2025. doi: 10.1038/s41592-025-02720-4.
- Liao, J., Zheng, D., Zhang, H., Zhang, L., Hu, M., and Bao, C. Fine-tuning AlphaFold with limited cryo-EM observations. *Communications Chemistry*, 9:95, 2026. doi: 10.1038/s42004-026-01899-7.
- Maddipatla, A., Sellam Bojan, N., Bojan, M., Masalitin, V., Vedula, S., Schanda, P., Marx, A., and Bronstein, A. M. Experiment-guided AlphaFold3 resolves measurement-consistent protein ensembles. *Nature biotechnology*, 2026. doi: 10.1038/s41587-026-03166-5.
- Mattingly, H. H., Evans, L., and Cossio, P. Measurement-limited learning of conformational heterogeneity in cryo-electron microscopy. *arXiv:2606.14449v1*, 2026. URL <https://arxiv.org/abs/2606.14449v1>.

- Narnhofer, D., Habring, A., Holler, M., and Pock, T. Posterior-Variance-Based Error Quantification for Inverse Problems in Imaging. *arXiv:2212.12499*, 2024. URL <https://arxiv.org/abs/2212.12499>. Archived manuscript dated 1 August 2024.
- Nguyen, N. D. and Pham, B. ARCHER: Amortized cross-specimen pose estimation for cryo-electron microscopy. *arXiv:2608.22029v1*, 2026. URL <https://arxiv.org/abs/2608.22029v1>.
- NIST. Digital Library of Mathematical Functions, Section 3.5: Quadrature, 2026. URL <https://dlmf.nist.gov/3.5>.
- Ortiz, S., Stanisic, L., Rodriguez, B. A., Rampp, M., Hummer, G., and Cossio, P. Validation tests for cryo-EM maps using an independent particle set. *Journal of structural biology: X*, 2020. doi: 10.1016/j.jsbx.2020.100032.
- Punjani, A. and Fleet, D. J. 3DFlex: determining structure and motion of flexible proteins from cryo-EM. *Nature Methods*, 2023. doi: 10.1038/s41592-023-01853-8.
- Qu, H., Wang, X., Zhang, G., Peng, J., and Chen, T. GEM: 3D Gaussian Splatting for Efficient and Accurate Cryo-EM Reconstruction. *arXiv:2509.25075*, 2025. URL <https://arxiv.org/abs/2509.25075>.
- Rangan, A. V., Tang, W.-S., Cossio, P., Zhang, K., and Grigorieff, N. Estimating the tails of the spectrum of the Hessian of the log-likelihood for ab-initio single-particle reconstruction in electron cryomicroscopy. *arXiv:2411.13263v2*, 2024. URL <https://arxiv.org/abs/2411.13263v2>.
- Sanchez-Garcia, R., Saur, M., Vargas, J., Poelking, C., and Deane, C. M. CESPED: a new benchmark for supervised particle pose estimation in Cryo-EM. *arXiv:2311.06194v2*, 2024. URL <https://arxiv.org/abs/2311.06194v2>.
- Sanchez-Garcia, R., Berndt, A., Apelbaum, A., Reeks, J., Williams, P. A., Poelking, C., Deane, C. M., and Saur, M. Supervised Deep Learning for Efficient Cryo-EM Image Alignment in Drug Discovery with cryoPARES. *bioRxiv, version 6*, 2026. doi: 10.1101/2025.03.04.641536. URL <https://www.biorxiv.org/content/10.1101/2025.03.04.641536v6>.
- Scheres, S. H. W. A Bayesian view on cryo-EM structure determination. *Journal of Molecular Biology*, 415:406–418, 2012. doi: 10.1016/j.jmb.2011.11.010.
- Scheres, S. H. W. and Chen, S. Prevention of overfitting in cryo-EM structure determination. *Nature Methods*, 9: 853–854, 2012. doi: 10.1038/nmeth.2115.
- Schwab, J., Kimanius, D., Burt, A., Dendooven, T., and Scheres, S. H. W. DynaMight: estimating molecular motions with improved reconstruction from cryo-EM images. *Nature Methods*, 2024. doi: 10.1038/s41592-024-02377-5.
- Shabeeb, Z., Saeedi, D., Tsui, D., Jamali, V., and Aghazadeh, A. cryoSENSE: Compressive sensing enables high-throughput microscopy with sparse and generative priors on the protein cryo-EM image manifold. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 34072–34083, 2026. URL https://openaccess.thecvf.com/content/CVPR2026/html/Shabeeb_cryoSENSE_Compressive_Sensing_Enables_High-throughput_Microscopy_with_Sparse_and_Generative_CVPR_2026_paper.html.
- Shekarforoush, S., Lindell, D. B., Brubaker, M. A., and Fleet, D. J. Reconstructing Heterogeneous Biomolecules via Hierarchical Gaussian Mixtures and Part Discovery. *arXiv:2506.09063*, 2025. URL <https://arxiv.org/abs/2506.09063>.
- Shihab, I. F., Akter, S., and Sharma, A. CalPro: Prior-aware evidential conformal prediction with structure-aware sensitivity bounds for protein structures. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, 2026. URL <https://proceedings.mlr.press/v306/shihab26b.html>.
- Sorzano, C. O. S., Jiménez-Moreno, A., Maluenda, D., Martínez, M., Ramírez-Aportela, E., Krieger, J., Melero, R., Cuervo, A., Conesa, J., Filipovic, J., Conesa, P., del Caño, L., Fonseca, Y. C., Jiménez-de la Morena, J., Losana, P., Sánchez-García, R., Strelak, D., Fernández-Giménez, E., de Isidro-Gómez, F. P., Herreros, D., Vilas, J. L., Marabini, R., and Carazo, J. M. On bias, variance, overfitting, gold standard and consensus in single-particle analysis by cryo-electron microscopy. *Acta Crystallographica Section D Structural Biology*, 2022. doi: 10.1107/S2059798322001978.
- Ullrich, K., van den Berg, R., Brubaker, M., Fleet, D., and Welling, M. Differentiable probabilistic models of scientific imaging with the Fourier slice theorem. In *Proceedings of the 35th Conference on Uncertainty in Artificial Intelligence*, volume 115 of *Proceedings of Machine Learning Research*, 2020. URL <https://proceedings.mlr.press/v115/ullrich20a.html>.
- Van, C. T. S., Reboul, C. F., Caesar, J. J. E., Meana-Pañeda, R., Lountos, G. T., Deme, J. C., Bryant, O. J., Johnson, S., Piczak, C. T., Valkov, E., Lea, S. M., and

- Elmlund, H. Probabilistic single-particle cryo-EM ab initio 3D reconstruction in SIMPLE. *Acta Crystallographica Section D Structural Biology*, 2025. doi: 10.1107/S2059798325005686.
- Vargas, J., Florea, A. V., Pellegrino, S., Burguet-Castell, J., Gómez-Pedrero, J. A., Arruda Bezerra, G., and Bui, K. H. Gold-standard local validation enhances interpretability of cryo-em maps. Zenodo version 3; attached manuscript filename v5, 17 June 2026, 2026.
- Vilas, J. L., Tagare, H. D., Vargas, J., Carazo, J. M., and Sorzano, C. O. S. Measuring local-directional resolution and local anisotropy in cryo-EM maps. *Nature Communications*, 2020. doi: 10.1038/s41467-019-13742-w.
- Xu, S., Balanov, A., Singer, A., and Bendory, T. Bayesian perspective for orientation determination in cryo-EM with application to structural heterogeneity analysis. *Acta crystallographica. Section D, Structural biology*, 2026. doi: 10.1107/S2059798326001415.
- Yamazaki, K., Wada, M., Wada, Y., Tokuhisa, A., Ohga, K., Nakagawa, A., Akinaga, Y., Sasakura, Y., Katoh, T., Matsumoto, S., and Okuno, Y. Digital twin of biomacromolecular thermodynamics in cryo-EM data. *Communications Chemistry*, 9:293, 2026. doi: 10.1038/s42004-026-02077-5.
- Ye, C. T., Li, S., King, T., and Monakhova, K. QUTCC: Quantile Uncertainty Training and Conformal Calibration for Imaging Inverse Problems. *arXiv:2507.14760v2*, 2026. URL <https://arxiv.org/abs/2507.14760v2>.
- Zhang, H., Zheng, D., Wu, Q., Yan, N., Peng, H., Hu, Q., Peng, Y., Yan, Z., Shi, Z., Bao, C., and Hu, M. CryoPROS: Correcting misalignment caused by preferred orientation using AI-generated auxiliary particles. *Nature Communications*, 16:4565, 2025a. doi: 10.1038/s41467-025-59797-w.
- Zhang, J., Zhou, S., Dai, H., Liu, X., Wang, P., Fan, Z., Pei, Y., and Yu, J. CryoFastAR: Fast Cryo-EM Ab initio Reconstruction Made Easy. *arXiv:2506.05864v1*, 2025b. URL <https://arxiv.org/abs/2506.05864v1>.
- Zhao, C., Xu, Y., and Liu, S. When Variance Is Not an Error Map: Calibrated Uncertainty for Radiative Gaussian Splatting in Sparse-View CT. *arXiv preprint arXiv:2607.13682v4*, 2026. URL <https://arxiv.org/abs/2607.13682v4>.
- Zhong, E. D., Bepler, T., Berger, B., and Davis, J. H. Cryo-DRGN: reconstruction of heterogeneous cryo-EM structures using neural networks. *Nature Methods*, 18:176–185, 2021a. doi: 10.1038/s41592-020-01049-4.
- Zhong, E. D., Lerer, A., Davis, J. H., and Berger, B. CryoDRGN2: Ab Initio Neural Reconstruction of 3D Protein Structures From Real Cryo-EM Images. In *ICCV*, 2021b. URL https://openaccess.thecvf.com/content/ICCV2021/html/Zhong_CryoDRGN2_Ab_Initio_Neural_Reconstruction_of_3D_Protein_Structures_From_ICCV_2021_paper.html.

A. Finite Representations and Local Pose Bounds

We use smooth conjugate pairs

$$\begin{aligned} F_c(k) &= \sum_j c_j g_j(k) + \bar{c}_j g_j(-k), \\ g_j(k) &= e^{-\|k - \mu_j\|^2 / (2\sigma_j^2)}. \end{aligned} \quad (13)$$

This enforces Hermitian symmetry. Hard cutoff kernels from the earlier reconstruction prototype are unsuitable for an unqualified Taylor bound; the uncertainty implementation uses full Gaussians. Their analytic whole-space Gram matrix provides density-energy coordinates $\delta = T\beta$ with $T^*GT = I$, so $\|\beta\| \leq B$ is a density-energy assumption rather than an arbitrary coefficient scale.

Parameterize local rotations by a Lie-algebra vector and translations by detector shifts. Normalize their joint vector so $\|u_i\| \leq 1$, with physical radii incorporated in the operator. With $D_{ia} = \partial A_i / \partial u_{ia}$ at the pilot pose,

$$\begin{aligned} A_i(u_i)(c_0 + \delta) &= A_i(0)(c_0 + \delta) + J_i u_i \\ &\quad + \sum_a u_{ia} D_{ia} \delta + e_i. \end{aligned} \quad (14)$$

If H_i uniformly bounds the pilot's second derivative along all allowed segments and K_i bounds the corresponding operator derivative, then $\|e_i\| \leq (H_i + BK_i)/2$. Define $T_i(w_i) = [D_{i1}^\top w_i, \dots, D_{iq}^\top w_i]$. The interaction contributes at most

$$B \sum_i \|T_i(w_i)\|_{\text{op}} \leq B \sum_i \|T_i(w_i)\|_F. \quad (15)$$

The implemented Frobenius relaxation is a sum of Euclidean group norms. It ignores some sharing of δ across particles and can be conservative. Charging the entire first-order interaction to an isotropic remainder was substantially worse in development.

The same construction operates in the full supported voxel space. For voxel position x , box length d , rotation radius a and shift radius s_t , set

$$L(q, x) = \frac{2\pi}{d} (a \|k\| \|x\| + s_t \|q\|), \quad (16)$$

$$H(q, x) = \frac{2\pi}{d} a^2 \|k\| \|x\|. \quad (17)$$

The absolute second derivative of its complex Fourier column is bounded uniformly by $|C_i(q)|[L(q, x)^2 + H(q, x)]$ before whitening. Summing this column bound against $|\rho_0(x)|$ gives the pilot bound; its frequency/voxel Frobenius norm gives the operator bound. Analytic first derivatives supply the full-density interaction. Thus the implemented joint certificate covers the stated ambient density ball and local pose balls simultaneously. Fixed CTFs and known whitened Gaussian noise remain assumptions; CTF error, window correlations and heterogeneous densities are not covered by assertion.

A.1. Curvature and shared-density post-audits

Retaining the quadratic pose terms replaces the second-order remainder with a uniform cubic one (Appendix C.4). It adds structured pilot and density Hessian terms, and can substantially reduce width. The density component can be tightened further: it is one shared error across particles, not one independent error per image. Let $h = \mathcal{A}^*w - \ell$ and collect blocks $H_0 = h$, $H_i = T_i(w)$ and $H_{n+i} = U_i(w)/2$, where U_i contains all ordered second-derivative adjoint columns. Their coefficients are 1, u_i , $u_i \otimes u_i$, each of norm at most one. For any positive fixed block scales d_j ,

$$\left\| \sum_j H_j v_j \right\| \leq \sqrt{\sum_j d_j} \|[H_j / \sqrt{d_j}]_j\|_{\text{op}}. \quad (18)$$

Multiplication by B bounds the combined density residual and interactions. The pilot and cubic remainder terms remain unchanged. Equation (18) is an elementary operator-norm inequality. We compare it with separate spectral and triangle bounds, use $d_j = \|H_j\|_F$ for nonzero blocks, and retain the smallest valid bound. This post-audit uses fixed weights and preserves coverage; it does not inherit an optimality certificate for the new width.

First and second derivative adjoints require only five first and fifteen distinct second Fourier moment fields per particle. Direct summation or batched nonuniform transforms evaluate these fields on a larger supported grid without its pixel-space derivative Grams. Physical voxel-volume scaling, the inference center and the target all use that audit grid. This verifies larger finite density spaces, not arbitrary unresolved continuous density.

B. Finite-Space Development Experiments

B.1. Representation and exact-dictionary controls

A noiseless projection audit with 256 particles per split and a radius-six frequency band found relative held-out errors of approximately 1.8%, 2.7% and 7.7% for the finest tested Gaussian dictionary. These residuals demonstrate why noise-only intervals do not cover representation error.

The bounded-model coverage benchmark has 48 settings: 12 seeds, broad versus preferred orientations, and rotation radii 0.5° or 2° . Each setting uses 96 particles, 40 coefficients, 20 random bounded signals and estimator-specific adversarial controls. We report exact Gaussian noise-marginal coverage and cross-check with 2,000 paired noise draws per signal. Shared designs and paired noise do not constitute independent voxel trials.

The exact-dictionary comparison is plotted in Appendix H. The structured certificate's smallest tested coverage was 0.991 at nominal 0.95; omitting the quadratic remainder

gave 0.932. Median width relative to the no-data bound was 0.557 for broad views at 0.5° , and approximately one for preferred views. This is a usefulness limitation, not a calibration success to conceal with pooled averages.

B.2. Ambient-space audit

The fixed-pose audit uses 128 real acquisition geometries per dataset, a 24^3 voxel grid, a spherical support of radius 0.35 times the field, and prescribed signal-to-noise power ratio 0.1. A pilot derived from real particles and each reference generator are normalized to unit voxel Euclidean norm. We set $B = 2$ before seeing coverage; the triangle inequality covers all such normalized generator/pilot pairs. This simulation convention does not estimate the density radius of experimental data.

The initial dictionary has 33 non-null Gaussian modes. Two targets, a central average and an axial contrast, use Gaussian averaging widths 0.03 and 0.07 times the field. Restricted intervals, the same weights with a full-grid audit, enriched weights, and the ambient matrix-free solution use identical estimands. The larger-grid model is deliberately able to express errors outside the initial dictionary.

Table 2 reports median widths across the two targets. The full-space solver reached its 0.5% objective-gap tolerance in 12 of 12 cases, taking 2.4–11.1 seconds per target on the Mac. The enrichment solver reached the same outer tolerance in 8 of 12 cases within 100 rounds; nonconverged results retain valid bounds but do not establish optimality. The solvers optimize the conservative sum-width, so their sharper reported widths need not be ordered exactly by their objective gaps.

Dictionary-only intervals can be narrow because they exclude admissible density errors. For the EMPIAR-10028 central broad average, analytic coverage on the independent voxel-map generator is approximately 4.2×10^{-5} . The same image weights, with the ambient bias bound, cover at the prescribed level over the ambient ball. On each estimator’s constructed ambient bias boundary, the audited intervals attain nominal 0.95 Gaussian coverage up to numerical tolerance. These boundary calculations check the theorem and implementation; they are not independent evidence of practical class calibration. Broader targets have substantially smaller relative widths than fine targets. The prescribed norm radius and fixed-pose assumption are essential to interpreting that result.

B.3. Joint ambient and nonlinear pose uncertainty

We combine the full supported voxel space with the nonlinear pose bound, using the broad axial contrast, 128 particles per geometry and the same prescribed density radius. Rotation radii are 0.1° , 0.5° and 2° , with a shift radius of 0.01

working pixels in their joint unit ball. Signal cases include random density/pose boundaries, estimator-directed adversaries and the independent voxel reference under a coherent rotation. Table 12 shows that a valid but conservative interval can still be uninformative at a large nuisance radius. Fixed-pose intervals reach coverage as low as 0.904 on these allowed perturbed cases. The nonlinear certificate’s smallest tested analytic coverage is 0.9997; its conservatism requires scrutiny rather than being interpreted as superior calibration. The earlier first-order optimization-budget comparison is retained in the appendix.

The completed curvature comparison spans three geometries, two targets and three pose radii. All 18 quadratic solves reach the 0.5% objective-gap tolerance. At 0.5° , their width/no-data ratios are 0.128–0.178, versus 0.423–0.532 at 2° . Figure 2 shows the full comparison, including probability of establishing the correct target sign. Central averages become detectable at the intermediate radius; the deposited-map axial contrasts remain too small. Their averaging standard deviation is 0.07 times the field (16.5–33.8 Å), so this does not establish atomic-scale utility.

For both targets at 0.5° and 2° , the shared-density post-audit reduces the same weights’ interval widths to 0.50–0.65 of their original values. Independent local nonlinear stress tests optimize pose errors while eliminating the density ball exactly. Four starts for each bias sign find feasible adversaries, recomputed in float64 after MPS search; they supply lower bounds on worst bias, not global solutions. Their comparisons with the upper bounds are retained in the repository.

A finer-feature development probe uses a 32^3 grid, frequency radius 12 and averaging width 0.03 times the field. At 0.5° on EMPIAR-10028, the first-order bound equals the no-data interval; retaining curvature reduces its relative width to 0.449. The completed comparison spans both targets and all three geometries: all twelve quadratic solves converge, with relative widths 0.166–0.231 at 0.1° and 0.442–0.477 at 0.5° . These remain finite-grid results.

B.4. Bootstrap tradeoff and finite-grid stability

A fixed-pose benchmark fits 33- and 123-mode Gaussian models and all 2,517 supported voxels. It uses 999 source-group bootstrap resamples, four prescribed ridge strengths and 1,000 independent Gaussian measurement-noise replications per setting. Percentile/basic intervals, sampling covariance, posterior covariance and a same-estimator ambient bias audit are recorded. The group-resampling plan is reused across noise draws; binomial error bars are conditional on that plan.

Weakly regularized full-voxel bootstrap intervals have center coverage 0.938, 0.956 and 0.958 across the three geometries,

Table 2. Fixed-pose ambient audit: median interval width relative to the no-data width across two density targets. All three methods after auditing use the same supported voxel class. The first column instead assumes the initial 33-dimensional dictionary is correct and is not an honest interval for that larger class. All entries are development results.

Dataset	Averaging width / field	Dictionary only	Same weights, audited	Enriched	Full matrix-free
10028	0.03	0.008	0.978	0.689	0.687
10028	0.07	0.007	0.510	0.053	0.050
10049	0.03	0.020	1.016	0.718	0.716
10049	0.07	0.020	0.636	0.075	0.073
10076	0.03	0.024	0.949	0.721	0.719
10076	0.07	0.023	0.462	0.072	0.072

with widths smaller than the robust audit. Strong regularization can cause substantial undercoverage: EMPIAR-10028 center coverage drops to approximately 0.339. A conservative bias audit therefore does not dominate bootstrap width in every setting. Figure 15 retains the regularization sweep, including cases favorable to bootstrap.

A fixed-pose refinement study audits 48 coarse-grid solutions on a common 64^3 supported space. Grids 16, 24, 32 and 48 preserve physical field, density-energy norm and noise units. Width ratios between fine audit and original calculation range from 0.72–1.45 at grid 16 and approximately 0.999–1.006 at grid 32. Nonlinear Fourier-moment audits additionally cover both targets, both pose radii and all three geometries on grids 24, 32 and 64. Discretizations are not identical nested classes; ratios need not increase monotonically. Agreement is a numerical sensitivity check, not a continuous-density guarantee.

C. Additional Derivations

C.1. An identifiability lower bound

Proposition 1. *Let two allowed parameters induce $N(\mu_0, I)$ and $N(\mu_1, I)$, with targets t_0, t_1 . Any interval covering each target with probability at least $1 - \alpha$ has expected length under the first law at least*

$$|t_1 - t_0|[1 - 2\alpha - \text{TV}(P_0, P_1)]_+, \quad (19)$$

where $\text{TV}(P_0, P_1) = 2\Phi(\|\mu_1 - \mu_0\|/2) - 1$.

Proof. Under P_0 , the probability of containing t_1 is at least $1 - \alpha - \text{TV}(P_0, P_1)$ by the definition of total variation. The union bound gives simultaneous containment of both targets with probability at least the displayed positive part. The length on that event is at least their separation. The total variation formula follows by projecting onto the line between the Gaussian means. $\square$

Two-point expected-length and between-class modulus bounds are classical; see Cai & Low (2004). This standard two-point argument also applies to different allowed

nuisance settings. Exact null directions give a nonvanishing uncertainty floor even with perfect image prediction. It does not assert that a finite numerical search finds the largest possible separation.

Corollary 1 (Fixed-length Gaussian testing bound). *Under $y \sim N(\mathcal{A}\rho, I)$ and $\|\rho - \rho_0\| \leq B$, any interval of deterministic length $2q$ covering every target $\langle \ell, \rho \rangle$ with probability at least $1 - \alpha$, $0 < \alpha < 1/2$, satisfies, for each nonzero $h \in H$,*

$$q \geq |\langle \ell, h \rangle| \min \left\{ \frac{B}{\|h\|}, \frac{z_{1-\alpha}}{\|\mathcal{A}h\|} \right\}, \quad (20)$$

where the second ratio is infinite when $\mathcal{A}h = 0$.

Proof. Take $\rho_{\pm} = \rho_0 \pm \lambda h$ with $\lambda \|h\| \leq B$ and $\lambda \|\mathcal{A}h\| < z_{1-\alpha}$. The observation laws have total variation less than $1 - 2\alpha$. If $2q < 2\lambda |\langle \ell, h \rangle|$, the two target-containment events are disjoint, whereas their respective coverage probabilities force total variation at least $1 - 2\alpha$, a contradiction. Taking the limit at the strict boundary proves the result. $\square$

This is a classical testing argument, not a new minimax theorem. We evaluate $h = \ell - \mathcal{A}^*w$ and $h = \ell$ and retain the larger lower bound. Analytic target integrals and quadrature error bounds evaluate the pairings and norm denominators. The latter use upper bounds, preserving conservatism in real arithmetic; roundoff remains unvalidated. The result concerns deterministic interval length, unlike the preceding expected-length statement. Nominal poses form a submodel of the bounded-pose class, but this lower bound alone cannot establish sharpness for the larger class.

The same argument gives a generic efficiency certificate, following the optimal-recovery connection (Donoho, 1994; Armstrong & Kolesár, 2018). Let $U = \mathcal{F}_H(w)$ and $L > 0$ be a feasible fixed-pose dual value, with $(U - L)/U \leq \varepsilon$. Scaling its dual vector by $\kappa = z_{1-\alpha}/z_{1-\alpha/2}$ makes it feasible for the testing bound. Thus the optimum deterministic half-length q^* obeys $q^* \geq \kappa L$, whereas the reported half-width satisfies $q \leq U$. Consequently

$$\frac{q}{q^*} \leq \frac{z_{1-\alpha/2}}{z_{1-\alpha}(1 - \varepsilon)}. \quad (21)$$




Figure 2. Curvature retained versus a first-order expansion with bounded second-order remainder, under the same density and pose classes. Dashed curves use the first-order procedure; solid curves retain curvature. The lower panels use the independent deposited-map generator under a coherent allowed rotation. Sign power and interval width expose limitations that high coverage alone conceals. These intervals precede the shared-density post-audit.

For $\alpha = 0.05$ and $\varepsilon = 0.005$, this is 1.198. The comparison applies to the continuous fixed-pose class; a post-hoc pose enlargement has no such gap certificate.

C.2. Gaussian density-energy Gram matrix

For isotropic paired features, let

$$P_{jk} = \left(\frac{2\pi\sigma_j^2\sigma_k^2}{\sigma_j^2 + \sigma_k^2} \right)^{3/2} \exp \left[-\frac{\|\mu_j - \mu_k\|^2}{2(\sigma_j^2 + \sigma_k^2)} \right], \quad (22)$$

$$Q_{jk} = P_{jk} \text{ with } \mu_j - \mu_k \text{ replaced by } \mu_j + \mu_k.$$

Integrating products of Gaussians and applying Parseval gives real and imaginary coefficient blocks $2(P + Q)$ and $2(P - Q)$, with zero cross block. A zero-centered imaginary coefficient is exactly null. Numerically excluded eigenmodes are reported; exclusion changes the stated Gaussian subspace. In an ambient audit they remain available to the larger space rather than being silently assigned zero uncertainty.

C.3. Uniform nonlinear remainder

For a Gaussian of width σ at radial distance r , bounds on value, gradient and Hessian are

$$g_0(r) = e^{-r^2/(2\sigma^2)}, \quad (23)$$

$$g_1(r) = rg_0(r)/\sigma^2, \quad (24)$$

$$g_2(r) = (1 + r^2/\sigma^2)g_0(r)/\sigma^2. \quad (25)$$

On an interval of possible distances, g_0 is largest at the lower endpoint; g_1 and g_2 attain their largest value at the point nearest σ . A rotation radius $a \leq \pi$ gives displacement at most $2\|k\|\sin(a/2)$, speed at most $a\|k\|$ and acceleration at most $a^2\|k\|$. A shift radius s gives phase speed at most $2\pi\|q\|s/\text{box}$.

The product rule bounds the second derivative by the sum of squared-speed times g_2 , acceleration times g_1 , twice speed times phase-speed times g_1 , and squared phase-speed times g_0 . Sum conjugate-pair bounds and apply the fixed CTF/noise scale. A coefficient-column Frobenius bound controls the operator term; the pilot-weighted bound controls its pilot term. Taylor's integral formula gives Equation (14). Evaluating a Hessian only at the nominal pose

would not establish this uniform result.

C.4. Retaining quadratic pose terms

The isotropic quadratic remainder can dominate an interval. A development extension retains the curvature and uses a cubic remainder instead. Put $E_{iab} = \partial^2 A_i / \partial u_{ia} \partial u_{ib}$ at zero, $Q_i(w)_{ab} = w_i^\top E_{iab} \rho_0$, and let $U_i(w)$ collect columns $E_{iab}^\top w_i$ over all ordered pairs. Since $\|u_i \otimes u_i\|_F \leq 1$, a valid bias bound is

$$\begin{aligned} b_2(w) = & B \|\ell - \mathcal{A}^* w\| + \sum_i \|J_i^\top w_i\| \\ & + B \sum_i \|T_i(w)\|_F + \frac{1}{2} \sum_i \|Q_i(w)\|_F \\ & + \frac{B}{2} \sum_i \|U_i(w)\|_F + \sum_i \gamma_{3i} \|w_i\|. \end{aligned} \quad (26)$$

For a Fourier voxel column, define $T(q, x) = 2\pi a^3 \|k\| \|x\| / d$. Its third derivative along any allowed pose segment is bounded by $|C|[L^3 + 3LH + T]$ before whitening. The three summands follow by differentiating the phase exponential three times. Pilot contraction and a column Frobenius bound give H_{3i} and K_{3i} , hence $\gamma_{3i} = (H_{3i} + BK_{3i})/6$. Taylor’s integral remainder and Cauchy–Schwarz then give Theorem 1 with b_2 . Sharing of density errors and rank-one quadratic pose constraints are still relaxed.

Analytic column Hessians and pixel-space Gram factors implement these Euclidean groups. Mixed finite differences and independent nonlinear projections test the implementation. The full three-geometry, two-target comparison is shown in Figure 2. The spectral post-audit follows Equation (18); its proof is the operator-norm bound applied to the concatenated vectors $[\sqrt{d_j} v_j]_j$, whose squared norm is at most $\sum_j d_j$.

C.5. An independent noise-scale upper bound

Known noise is not necessary if an independently calibrated upper bound is available. A standard construction assumes independent calibration coordinates $z_j = \mu_j + \sigma e_j$, $e_j \sim N(0, 1)$, with arbitrary fixed means and independent inference noise. For ν coordinates, set

$$\bar{\sigma} = \|z\| / \sqrt{\chi_{\nu, \beta}^2}. \quad (27)$$

The probability that $\bar{\sigma} \geq \sigma$ is at least $1 - \beta$: the central chi-square pivot is exact at zero mean, and a noncentral chi-square has a smaller lower tail. Conditional on calibration, weights may use $\bar{\sigma}$. On this event their inference standard deviation is at most $s_+ = \bar{\sigma} \|w\|$. For $\alpha < 1/2$, $q(s_+, b, \alpha) \geq b$, and Gaussian coverage decreases with the actual noise standard deviation up to s_+ .

Integrating the conditional bound yields coverage at least $(1 - \alpha)(1 - \beta) \geq 1 - \alpha - \beta$.

This extension does not justify independence of windowed Fourier pixels or estimate arbitrary covariance. With 128 calibration coordinates and $\beta = 0.01$, a 20,000-draw development control gives scale-bound failure 0.0103 under independence. Equicorrelation 0.5 raises that failure to 0.503. Such a calibration set violates the premise even if a particular resulting interval remains conservative. Experimental whitening remains unresolved.

C.6. Simultaneous targets and sensitivity budgets

For K targets and parameter-budget choices fixed independently of inference noise, applying Theorem 1 with α/K gives simultaneous coverage by the union bound. Correlated Gaussian maxima can reduce the noise component if the joint weight covariance is accounted for. Selecting the most favorable point on a sensitivity surface after inspecting inference data requires such simultaneous control. A density radius supplied by simulation truth is an oracle control and cannot be reported as a practical calibration method.

D. Survey and Experimental Audit Trail

The repository’s living survey distinguishes density uncertainty, pose uncertainty, conformational variability/populations, map-processing confidence and atomic-model uncertainty. Its broad Europe PMC retrieval ledger contains 16,969 deduplicated records, mostly application papers; that count is not a count of reviewed methods. The initial survey checkpoint contained 89 candidate works; the current ledger contains 101, with local full-text availability recorded separately. Recent items include LocScale-2.0, CryoDiff, CryoSplat, GEM, measurement-limited heterogeneity, population-counting work and the CAHRA dataset-construction preprint. Preprints and final publications are distinguished, and titles/versions are checked against primary sources.

The full experimental protocol draft records unimplemented comparisons and required failure controls. Fixed-pose resampling must be distinguished from rerunning alignment; a mathematical diagonal-Gaussian control must be distinguished from reproducing Ullrich et al.’s full software; cryoDRGN backprojection must be distinguished from training a neural model. Exposure-group splits do not undo the dependence introduced by full-data consensus poses. All development results, including the initially incomplete covariance-factor benchmark, remain identifiable by filename and status.

E. Continuous Fourier Audit

The sinc identity follows by integrating the two conjugate expansions of a real Fourier field. For a target component with center μ and width σ_ℓ , the one-dimensional truncated transform is

$$e^{-2\pi i k \mu - 2\pi^2 \sigma_\ell^2 k^2} \times [\Phi((1/2 - \mu)/\sigma_\ell + 2\pi i \sigma_\ell k) - \Phi((-1/2 - \mu)/\sigma_\ell + 2\pi i \sigma_\ell k)]. \quad (28)$$

Products over coordinates and sums over target components give $\mathcal{A}\ell$. Gaussian-product identities give $\|\ell\|^2$. Independent tensor quadrature verifies these analytic expressions, while an independent conic problem formed from the joint Hilbert Gram of $[\ell, \mathcal{A}_1^*, \dots, \mathcal{A}_m^*]$ checks optimization and dual feasibility.

For r Gauss–Legendre nodes on a unit interval, the classical error constant is $C_r = (r!)^4 / [(2r+1)((2r)!)^3]$ (NIST, 2026). Put $M_a = \max_j |k_{ja}|$ and $E_r = \sum_a C_r (4\pi M_a)^{2r}$. Positive tensor weights ω_t satisfy

$$\left| \|f_w\|^2 - \sum_t \omega_t f_w(x_t)^2 \right| \leq E_r \left(\sum_j |c_j| \right)^2. \quad (29)$$

To verify the bound, expand f_w^2 into exponentials with coordinate frequencies at most $2M_a$. Apply the one-dimensional derivative remainder and telescope the three integration/quadrature operators. Their positive weights have total mass one, and the expansion’s coefficient absolute sum is bounded by $(\sum_j |c_j|)^2$.

Similarly, the error in the continuous Gram action obeys

$$\|(G - G_Q)w\| \leq \sqrt{2} E_r \|C/\sigma\|_2 \sum_j |c_j|. \quad (30)$$

Adding the scalar remainder to the squared primal residual and the vector remainder to the dual feasibility denominator preserves the continuous objective bracket in real arithmetic. Two type-3 NUFFTs apply G_Q . At frequency radius five, quadrature order 40 makes the analytic kernel remainder at most 3.21×10^{-23} before coefficient scaling. This quadrature theorem does not certify FINUFFT’s floating-point error; requested tolerance 10^{-12} is checked numerically against independent sinc Grams.

An orthonormal constant-cell coefficient is scaled by the square root of inverse cell volume. Its exact transform includes $\prod_a \text{sinc}(k_a/d)$ and a half-cell phase relative to the centered FFT grid. This prevents a point-sample approximation from being mislabeled exact continuous physics. For an orthogonal cell projection P , $\|h\|^2 = \|Ph\|^2 + \|(I - P)h\|^2$. The omitted component is the continuous analogue of the omitted-dictionary term, and can affect coverage even when successive grid widths look similar.

E.1. Continuous nonlinear-pose bound

Proposition 2. Let $H = L^2([-1/2, 1/2]^3)$, $\|\rho - \rho_0\| \leq B$, $P = \|\rho_0\|$ and $\|u_i\| \leq 1$. Put $h = A_0^* w - \ell$. Suppose L_{12} bounds the continuous norm of the sum of first and half-scaled second pose-adjoint fields. For rotation radius a and shift radius s_t in field units, define $R = \sqrt{3}/2$ and

$$L_{iq} = 2\pi(a \|k_{iq}\| R + s_t \|q\|), \quad (31)$$

$$H_{iq} = 2\pi a^2 \|k_{iq}\| R, \quad T_{iq} = 2\pi a^3 \|k_{iq}\| R, \quad (32)$$

$$K_{3i}^2 = \sum_q |C_{iq}/\sigma|^2 (L_{iq}^3 + 3L_{iq}H_{iq} + T_{iq})^2. \quad (33)$$

Then Equation (11) bounds the scalar bias for every density and pose in the class. Combined with Theorem 1, it yields a continuous-density interval under the same independent Gaussian-noise assumptions.

Proof. The exact bias splits as $\langle h, \delta \rangle + \langle (A(u)^* - A_0^*)w, \rho_0 + \delta \rangle$. The first term is at most $B \|h\|$. The second pairs a Taylor polynomial plus remainder with a density of norm at most $P + B$. Along a rotation segment, the first three frequency-derivative norms are bounded by $a \|k\|$, $a^2 \|k\|$ and $a^3 \|k\|$; the shift phase is linear. The third exponential derivative is bounded by $L^3 + 3LH + T$. The cube has unit volume, so summing squared column bounds gives the realified Hilbert–Schmidt bound K_{3i} . Taylor’s integral remainder contributes $1/6$; Cauchy–Schwarz with w_i completes the bound. $\square$

To compute L_{12} , use orthonormal symmetric vectorization of the Hessian: diagonal entries are unchanged and off-diagonals are multiplied by $\sqrt{2}$. Its coefficient vector has norm $\|u_i\|^2 \leq 1$, so fifteen columns replace twenty-five ordered columns exactly. The derivative fields are polynomials of degree at most two times Fourier exponentials. Their Gram products have coordinate degrees at most four.

For quadrature order r , $n = 2r$, and $\omega_a = 4\pi \max_j |k_{ja}|$, define

$$D_\beta(\omega) = \sum_{l=0}^{\min(\beta, n)} \binom{n}{l} \frac{\beta!}{(\beta-l)!} 2^{l-\beta} \omega^{n-l}, \quad (34)$$

$$E_{\text{poly}} = C_r \sum_a \max_{0 \leq \beta \leq 4} D_\beta(\omega_a). \quad (35)$$

Leibniz differentiation, the positive Gauss remainder kernel and tensor telescoping bound each monomial-exponential integration error by E_{poly} . If column j has absolute polynomial-Fourier coefficient sum at most M_j , the Gram-entry error is at most $E_{\text{poly}} M_j M_l$ and its operator norm at most $E_{\text{poly}} \sum_j M_j^2$. After dividing each column and its M_j by $\sqrt{d_{\text{block}(j)}}$, add that error to the largest quadrature-Gram eigenvalue in Equation (18). Separate and joint spectral bounds, and a padded Frobenius triangle bound, each remain valid; their minimum supplies L_{12} .

This integrates known adjoint fields, not the unknown density. Direct nonlinear Fourier fields check derivatives and cubic remainders independently. Exact constant-cell projections test reference generators. At a fixed nonzero pose, sinc integrals give the feasible worst bias over the entire continuous density ball, $|\langle \rho_0, (A(u)^* - A_0^*)w \rangle| + B \|A(u)^*w - \ell\|$; it is a lower bound on the global pose supremum. Floating-point and NUFFT tolerances remain numerical checks rather than validated arithmetic.

E.2. Integrated-moment cubic remainder

The earlier K_{3i} uses the maximum spatial radius everywhere. A sharper scalar-estimator bound integrates known Fourier columns over the cube. Put

$$c_2 = 12^{-1/2}, \quad c_4 = (3/12^2)^{1/4}, \quad c_6 = (15/12^3)^{1/6}. \quad (36)$$

For any $v \in \mathbb{R}^3$, $\|v^\top x\|_{L^p(D)} \leq c_p \|v\|$ for $p = 2, 4, 6$. To see this, expand the even moment under unit-volume cube integration. Odd-coordinate terms vanish, and each remaining term is nonnegative. The coordinate moments $1/12$, $1/80$ and $1/448$ are bounded by those of a centered Gaussian with variance $1/12$. Termwise comparison gives the stated constants. This is an integration identity for the operator, not a Gaussian model of the density.

Proposition 3 (Integrated scalar remainder). *Under Proposition 2's density and pose classes, define*

$$L_p(k, q) = 2\pi(a \|k\| c_p + s_t \|q\|), \quad (37)$$

$$H_p(k) = 2\pi a^2 \|k\| c_p, \quad T_2(k) = 2\pi a^3 \|k\| c_2, \quad (38)$$

$$D(k, q) = L_6(k, q)^3 + 3L_4(k, q)H_4(k) + T_2(k). \quad (39)$$

For $c_{iq} = (w_{R,iq} + iw_{I,iq})C_{iq}/\sigma$, the uniform scalar cubic remainder is bounded by

$$r_{\text{mom}}(w) = \frac{B+P}{6} \sum_{i,q} |c_{iq}| D(k_{iq}, q). \quad (40)$$

It can replace the cubic term in Equation (11), preserving the same conditional coverage guarantee.

Proof. Along a normalized pose segment, the phase $\phi(t, x)$ has derivative norms $\|\phi'\|_p \leq L_p$, $\|\phi''\|_p \leq H_p$ and $\|\phi'''\|_2 \leq T_2$. This follows from $\|k^{(j)}(t)\| \leq a^j \|k\|$ and linear detector translation. The exact third derivative of $e^{i\phi}$ is $e^{i\phi}(i\phi''' - 3\phi'\phi'' - i(\phi')^3)$. Minkowski and Hölder give its L^2 norm at most D . Triangle inequality with the fixed coefficients c_{iq} , Taylor's integral weight $1/6$, and Cauchy-Schwarz against a density of norm at most $B+P$ prove Equation (40). Moreover, $\sum_q |c_{iq}| D_{iq} \leq \|w_i\| (\sum_q |C_{iq}/\sigma|^2 D_{iq}^2)^{1/2}$, so retaining actual weights is no worse than this integrated Hilbert-Schmidt bound. $\square$

The refinement applies classical inequalities to the Fourier-slice operator. Independent numerical integration checks the moment envelope, matrix-exponential derivatives, and exact nonlinear remainder. Those tests complement the proof; they are not exhaustive verification or validated arithmetic. The first development evaluation uses all original weights without refitting. A subsequent frozen protocol reserves different acquisition subsets and new signal/noise seeds for the one- and two-degree settings.

F. Development After Independent Review

The original frozen protocols and their outcomes are unchanged. This section reports subsequent development, including negative experimental sensitivity results. It does not establish that the current method meets the usefulness requirements raised by the independent review. The old frozen evaluators abort on detected undercoverage before retaining the offending record. Neither completed study encountered that branch, but the failure policy was flawed. Their source remains reproducible; new diagnostic evaluators persist results before flagging an audit failure.

F.1. Matrix-free shared fields and a numerical failure budget

Each first or second pose-derivative field is a Fourier sum multiplied by one of ten monomials $1, x, y, z, x^2, y^2, z^2, xy, xz, yz$. Summing coefficient combinations before transforming gives a field/adjoint pair with ten type-3 NUFFTs per direction. No global particle-by-particle field Gram is formed. Positive group scales d_j may come from streamed block norms or from polynomial coefficient envelopes. They need not equal exact block norms. Identically zero blocks may be omitted; blocks that can become nonzero during optimization require positive scales. Particle-specific pose radii are accepted.

An approximate leading eigenvalue is insufficient for an upper audit. Random starts have a long history in spectral estimation (Kuczynski & Wozniakowski, 1992); the following elementary construction states our particular numerical probability budget. For a fixed positive semidefinite $C = FF^\top$, draw m independent standard Gaussian probes g_j , independently of the estimator-selection process. Set $c = \sqrt{2} \operatorname{erf}^{-1}(\delta^{1/m})$. Fix a leading unit eigenvector v before the draws. Since

$$\Pr\{\max_j |v^\top g_j| < c\} = \delta,$$

with probability at least $1 - \delta$, simultaneously for every integer $p \geq 1$,

$$\lambda_{\max}(C) \leq \left(\frac{\max_j \|C^p g_j\|}{c} \right)^{1/p}. \quad (41)$$

The proof retains the leading eigenvector component of $C^p g_j$ in its norm. Stable normalization tracks the product of norms in log space. The minimum across powers is valid on the same event; a deterministic trace bound may cap it. A Rayleigh value is recorded only as a lower diagnostic.

Add the analytic integration pad to this spectral upper value before using the shared-field inequality. Allocate $\alpha_{\text{noise}} + \delta \leq \alpha$. A union bound then includes numerical randomization in the stated failure budget. Candidate selection using the same probes is not licensed: weights and scales must be fixed before fresh final probes, or all candidates need a simultaneous budget. Operator roundoff and NUFFT approximation remain numerical checks, not validated arithmetic.

Completed radius-five audits with 128 and 1,024 particles take 50.3 and 353.3 seconds under concurrent Mac load, with measured peak resident memory of 390 and 434 MB. Their relative half-widths are 0.438 and 0.740, respectively. They use different fitted estimators, so the latter result does not measure a particle-count benefit. The 128-particle randomized field bound is 1.157 times the dense comparator. A distinct 10,000-particle, radius-12 audit with 220 independent Fourier pairs per particle completes in 4,965 seconds with 1.95 GB peak resident memory. Its randomized eigenvalue upper is 1.136 times the Rayleigh lower diagnostic. Relative half-width is 0.655 at 1° and $0.5 \text{ \AA}$ translation. This uses deterministic matched-filter weights for a broad target, not optimized reconstruction weights. It establishes numerical scale under supplied simulation bounds, not useful fine-feature inference.

F.2. Including pose uncertainty in estimator optimization

Fix the group scales, let $F(w)$ be the scaled quadrature field matrix, and let $a_q(w)$ be each real/imaginary weight pair's Euclidean amplitude. A nonnegative matrix M maps these amplitudes to polynomial coefficient envelopes. With integration constant E , $R = B + P$, and integrated cubic coefficients $r_q \geq 0$, the revised sum objective is

$$J(w) = z \|w\| + B \|\ell - \mathcal{A}^* w\| + p(w) + \sum_q r_q a_q(w),$$

$$p(w) = R \sqrt{\sum_j d_j \sqrt{\|F(w)\|_{\text{op}}^2 + E \|M a(w)\|^2}}. \quad (42)$$

This is convex: F is linear, M is nonnegative, and the Euclidean norm is monotone on the nonnegative orthant. It is an upper objective for the prescribed class, not the exact minimax nonlinear-pose risk. Approximate singular modes and log-sum-exp smoothing guide numerical optimization. The final primal upper audit uses Equation (41) with fresh probes.

A feasible dual lower bound remains available without an optimizer-success claim. Choose $v \in H$ and supporting linear functionals g_p, g_r such that

$$\|v\| \leq B, \quad \|\mathcal{A}v - g_p - g_r\| \leq z,$$

$$g_p^\top x \leq p(x), \quad g_r^\top x \leq \sum_q r_q a_q(x) \quad \text{for all } x.$$

The four support inequalities in Equation (42) yield $J(x) \geq \langle v, \ell \rangle$ by cancellation. Unit singular-direction outer products and their convex mixtures have nuclear norm at most one, giving valid spectral supports even when the selected modes are not leading. Coefficient and cubic norm supports are combined with them. A padded continuous Gram-action error bounds the remaining dual constraint; scaling all three dual quantities by $\min(1, z/D)$ for an upper bound D makes it feasible. The implementation details and full support argument are recorded with the source.

An independently assembled small CLARABEL conic instance brackets its optimum with relative gap 2.60×10^{-4} . The first larger, 20-iteration development solve reduces a one-degree interval to 0.255 of no data but has a 97.8% gap; it is not converged. Six subsequent feasible nonlinear searches reach 0.432 of its bias upper bound, and minimum reference sign power is below 10^{-100} . The first completed 100-iteration case in the three-stack grid, 10049 at 0.5° and $0.5 \text{ \AA}$ translation, has relative width 0.221 and a 25.3% gap. Feasible searches reach 0.608 of its bias upper, but minimum reference sign power is only 8.44×10^{-5} . Its completed one-degree counterpart has width 0.336, an 85.3% gap and minimum reference power 7.07×10^{-25} . A feasible convex mixture of twelve spectral dual supports reduces the gap only to 85.1%, without changing the estimator or interval.

The original 10028 half-degree run failed when only seven of eight requested eigenvectors converged; the failure and saved weights are retained. Checkpointed continuations use an isolated two-thread CPU runtime with unchanged block scales and fresh final spectral probes. Available partial modes can guide the approximate optimizer, but are never used as upper bounds. The completed 10076 half-degree continuation has width 0.161, a 5.0% gap, minimum reference sign power 0.670, and feasible stress bias reaching 0.650 of the upper bound. This is a broad-target conditional result; the remaining grid is in progress. These intermediate outcomes do not establish broad practical usefulness.

Nonsmooth optimization and dual diagnostics. The longer legacy one-degree smooth probe fails before its final audit when only seven of eight modes converge; it has no certified result. To separate optimizer failure from estimator quality, we also restrict weights to $w = Vb$ and exchange spectral supporting cuts. For unit u, v , $g^\top w = v^\top F(w)u \leq \|F(w)\|_{\text{op}}$. Noise, density-residual and cubic pair norms

then form a small conic problem with constraints $t \geq g_j^\top Vb$. The unknown density class remains continuous. The restricted cut value is an optimization device, not a full-space lower bound; integration padding and fresh spectral probes remain in the final audit.

Conic cut multipliers obey $\mu_j \geq 0$ and $\sum_j \mu_j \leq R\sqrt{\sum d}$, so $\sum_j \mu_j g_j$ is a valid full-space pose dual support. Cubic cone duals, projected onto their known pair balls, avoid normalizing nearly zero weight pairs. More generally, minimizing the dual constraint defect over a convex hull of pose supports plus these pair balls can improve a lower bound without changing any interval. An independent small full spectral-norm conic problem checks the cut solver and support inequalities, including a nonzero optimum and an active zero pair.

On 10049 at two degrees and 0.5 Å, the original completed fit has relative width 0.761 and a 99.3% gap. Joint dual postprocessing reduces that gap to 63.1% with unchanged weights. Two consecutive thirty-round restricted-cut probes reduce width to 0.692 and 0.689. Their final local feasible stress biases reach 0.459 and 0.466 of the respective upper bounds; reference sign power remains below 10^{-240} . The follow-up's own joint dual gap is still 93.9%. These limited subspaces and dual directions do not establish full-space optimality or useful detection at two degrees.

A six-round conic trial optimizes all observation weights, projecting only the inner density residual onto a target-anchored continuous Hilbert basis. The projected residual is a lower relaxation, while the final audit retains the full density class. At rank 256 the new candidates have much larger ambient objectives than the input; the retained weights are unchanged. A fresh audit gives width 0.690 and gap 76.2%. A thirty-round adaptive continuous residual-enrichment follow-up takes 721 seconds and raises the valid lower bound to 9.200, reducing the original sum-objective gap to 18.2%, again with unchanged weights. This gap concerns the original triangle objective, not the subsequent tighter cross-term bound. These diagnostics distinguish improved lower bounds from an improved reconstruction.

A sixty-round full-weight continuation, with the latest eight candidates averaged before independent evaluation, changes the retained weights. It takes 2,330 seconds, gives a sum-objective upper bound 10.523 and feasible lower bound 9.845 (6.44% gap), and reduces the relative width to 0.646. The strongest saved feasible nonlinear bias reaches 0.521 of its upper bound. This establishes progress on the original objective, not useful reference sign detection or optimality of the subsequent joint expression.

A further sixty rounds use the sharp cubic constants and therefore reset the lower bound for the changed objective. They take 4,666 seconds and give upper/lower values

9.8961/9.5935 (3.06% gap), with relative width 0.607. The subsequent cross-term and known-pilot audits reduce that width to 0.548 and 0.503, respectively. The strongest saved nonlinear bias is 0.630 of the last upper, but minimum reference sign power remains below 10^{-118} . Those later widths do not inherit the sum-objective optimization gap.

Retaining the density/pose cross term. Let $h = \ell - A_0^*w$ and write the existing nonlinear adjoint perturbation as $D = Fv + e$, where $\|v\| \leq L = \sqrt{\sum d}$ and $\|e\| \leq r$. Here each particle lies in the joint scaled five-dimensional unit ball. Simultaneous maximal rotation and translation are excluded. For their product set, replace L^2 by $2 \sum_i d_i^{(1)} + 4 \sum_i d_i^{(2)}$ and scale f by the corresponding radius ratio. The separate phase bounds in the cubic remainder already cover either set. The two nuisance classes are reported separately. Suppose f bounds $\|Fv\|$ and c bounds $\|F^*h\|$. The scalar bias for $\rho = \rho_0 + u$ is $\langle u, h - D \rangle - \langle \rho_0, D \rangle$. Thus

$$b_{\text{joint}} = B\sqrt{\|h\|^2 + 2Lc + f^2} + Pf + (B + P)r \quad (43)$$

is an upper bound: expand $\|h - Fv\|^2$, bound the cross term by Lc , and apply Cauchy-Schwarz to u and ρ_0 . Take the minimum with the preceding triangle bound on the same spectral event. This elementary inequality is a post-audit refinement, not a new general theorem or a claim that the expression is convex in w .

Gaussian target moments evaluate $\langle \ell, F_j \rangle$ analytically on the cube, including boundary terms. For $I_j = \int_{-1/2}^{1/2} x^j g_{\mu,s}(x) e^{i\omega x} dx$, integration by parts gives

$$I_1 = (\mu + is^2\omega)I_0 - s^2[g_{\mu,s}(x)e^{i\omega x}]_{-1/2}^{1/2},$$

$$I_2 = (\mu + is^2\omega)I_1 + s^2I_0 - s^2[xg_{\mu,s}(x)e^{i\omega x}]_{-1/2}^{1/2}.$$

For the finite Fourier-polynomial product $\langle A_0^*w, F_j \rangle$, the existing quadrature error gives a cross-vector norm pad $ET\|m\|$, where T is the absolute Fourier coefficient sum and m contains the scaled column coefficient bounds. Ordinary numerical roundoff remains subject to the preceding caveat.

Table 3 retains the completed same-weight audits at this revision cutoff; the machine-readable summary lists every attempt and its source certificate seed. Each row uses its own source spectral event, without selecting a minimum over repeated random certificates. The 10076 half-degree conditional reference power increases from 0.670 to 0.968; the two-degree cut fit improves from width 0.689 to 0.598, but its reference power remains negligible. The saved nonlinear feasible stresses do not violate any new bound. This is evidence of reduced conservatism on these cases, not experimental density coverage.

An optional further constant refinement uses the sharp cube directional moments $\sup_{\|u\|=1} \mathbb{E}(u^\top X)^4 = 13/720$ and

Stack / fit	Target	Degrees	Old width	Joint width	Stress / upper	Min. power
10049	center	0.5	0.221	0.196	0.700	0.008
10049	center	1	0.336	0.299	0.527	1.0e-16
10028	center	0.5	0.146	0.125	0.754	1.000
10028	center	1	0.230	0.196	0.523	2.0e-20
10028	center	2	0.602	0.547	0.217	0 [†]
10049	center	2	0.761	0.701	0.251	7.4e-196
10049	contrast	0.5	0.253	0.227	0.741	7.0e-33
10049	contrast	1	0.372	0.337	0.574	3.8e-87
10049	contrast	2	0.840	0.775	0.271	0 [†]
10076	center	0.5	0.161	0.140	0.767	0.968
10076	center	1	0.243	0.209	0.532	0.002
10076	center	2	0.571	0.516	0.295	3.4e-128
10076	contrast	0.5	0.195	0.168	0.768	2.1e-24
10076	contrast	1	0.298	0.254	0.537	6.1e-58
10076	contrast	2	0.702	0.632	0.259	0 [†]
10049 / full	center	2	0.646	0.581	0.582	1.7e-175
10049 / cuts	center	2	0.689	0.598	0.539	1.6e-168
10049 / sharp fit	center	2	0.607	0.548	0.576	3.1e-151

Table 3. Same-weight residual/pose cross-term audits. Widths are relative to the no-data half-width. All cases use a broad $\sigma_\ell = 0.07$ target and a joint scaled pose ball with translation radius $0.5 \text{ \AA}$; simultaneous maximum rotation and translation are excluded. Each source cubic remainder and spectral event is retained. Stress ratios reuse saved feasible nonlinear poses; power uses three conditional reference-map scenarios ([†]: numerical underflow). These are exploratory numerical results, not experimental coverage labels.

$\sup_{\|u\|=1} \mathbb{E}(u^\top X)^6 = 205/36288$, for uniform cube X . Writing $p_j = u_j^2$, the sixth moment is $5/576 - \sum p_j^2/96 + \sum p_j^3/252$; pairwise averaging maximizes it at $p_j = 1/3$. Replacing the earlier Gaussian moment majorants in the same integrated remainder reduces the two-degree cut-fit width further to 0.572. For the full-weight averaged fit, the joint and sharp-moment refinements give widths 0.581 and 0.557, with feasible-over-upper ratios 0.582 and 0.608. Expanding to the product pose set increases the latter width to 0.616; saved joint-ball adversaries remain feasible examples but do not explore its extra poses. Minimum reference sign power remains below 10^{-155} even for the sharp joint-ball result. Independent tests exercise archive-scale Gaussian moments, the analytic cross-product quadrature pad, and the distinction between joint and product pose lifts. The original fit and every previous bound are retained. Neither refinement establishes useful fine-feature detection at one or two degrees.

Using the known pilot. The pilot term need not pay its full norm bound. If $p \geq \|F^* \rho_0\|$, Cauchy–Schwarz also gives $|\langle \rho_0, Fv \rangle| \leq Lp$. Hence the same spectral event supports

$$b_{\text{pilot}} = B\sqrt{\|h\|^2 + 2Lc + f^2} + \min(Pf, Lp) + (B+P)r. \quad (44)$$

The triangle alternative replaces the square root by $\|h\| + f$; their minimum remains valid. This changes neither the

unknown density class nor the weights. It is an elementary refinement of the known-function term, not a new general coverage theorem or an optimization-gap certificate.

For the known orthonormal constant-cell pilot on an m^3 grid, write $x = c + z$ inside each cell. For a uniform offset $z \in [-1/(2m), 1/(2m)]$, with $\text{sinc}(u) = \sin(\pi u)/(\pi u)$,

$$\begin{aligned} \mathbb{E}e^{2\pi i k z} &= \text{sinc}(k/m), \\ \mathbb{E}ze^{2\pi i k z} &= \frac{\text{sinc}'(k/m)}{2\pi i m}, \\ \mathbb{E}z^2e^{2\pi i k z} &= -\frac{\text{sinc}''(k/m)}{4\pi^2 m^2}. \end{aligned}$$

Expanding $x^\beta = (c + z)^\beta$ gives all ten degree-at-most-two Fourier moments from ten transforms of center-weighted cell coefficients. The orthonormal basis supplies $m^{-3/2}$ and the centered Fourier grid requires phase $\exp(i\pi \sum_a k_a/m)$. Thus the known discontinuous pilot is integrated over its cells, without replacing the unknown density by a cell representation. Near zero we evaluate the sinc derivatives by their Taylor series.

Independent composite cell quadrature and nonlinear finite differences check normalization and pose signs. Direct exponential sums over the actual 24^3 centers check all twenty lifted columns for three selected particles in each of the first ten archived audits. The largest relative pairing discrepancy is 3.49×10^{-13} ; this selected-column check is not a global floating-point error bound. The analytic formula

has no quadrature error, but its NUFFT and floating-point evaluation remain numerically checked rather than interval-validated.

Table 4 reports this separate same-weight variant. For the full-weight two-degree fit, the width falls from 0.557 to 0.510, and the saved feasible-over-upper ratio rises to 0.667. Its minimum reference sign power is still below 5×10^{-122} . The 10076 broad central feature has power 0.997 at half a degree and 0.214 at one degree; its two-degree width is 0.427, with negligible power. These conditional improvements leave fine-scale usefulness and experimental calibration open.

F.3. Constructive ambiguity from two feasible pose designs

Wide upper bounds alone cannot distinguish intrinsic ambiguity from a loose audit. We therefore construct lower bounds from two allowed pose designs, using the classical Gaussian testing argument (Donoho, 1994). This is a diagnostic on the same declared class, not a new minimax principle.

Let A_0, A_1 be two allowed observation operators, and let ρ_j belong to $\mathcal{C} = \{\rho : \|\rho - \rho_0^{\text{pilot}}\| \leq B\}$. Both designs have identity observation covariance. If

$$D = \|A_1\rho_1 - A_0\rho_0\| < 2\Phi^{-1}(1 - \alpha), \quad (45)$$

then every uniformly $1 - \alpha$ valid interval whose length is deterministic must have half-length at least $|\langle \ell, \rho_1 - \rho_0 \rangle|/2$. Indeed, if a shorter interval contained either feature with probability at least $1 - \alpha$, its membership test for one feature would distinguish the two hypotheses with both error probabilities at most α . The minimum sum of errors for their Gaussian distributions is $2\Phi(-D/2) > 2\alpha$, a contradiction. This conclusion does not bound the expected length of arbitrary variable-length intervals.

Write $p = A_1\rho_0^{\text{pilot}} - A_0\rho_0^{\text{pilot}}$ and choose a distance budget $\tau < 2\Phi^{-1}(1 - \alpha)$. For any observation weights w ,

$$G(w) = B \sum_{j=0}^1 \|\ell - A_j^*w\| - w^\top p + \tau \|w\| \quad (46)$$

upper-bounds the largest feasible feature gap for this fixed pair of pose designs. To see this, insert A_j^*w into the two feature pairings, apply the density-ball bounds, and use $w^\top(A_1\rho_1 - A_0\rho_0) \leq \tau \|w\|$.

The same weights yield explicit continuous witnesses. Put $h_j = \ell - A_j^*w$, choose $r_j \geq \|h_j\|$, and set

$$\tilde{\rho}_1 = \rho_0^{\text{pilot}} + Bh_1/r_1, \quad \tilde{\rho}_0 = \rho_0^{\text{pilot}} - Bh_0/r_0. \quad (47)$$

If $r_j = 0$, its perturbation is set to zero. Their mean difference is $p + \sum_j (B/r_j)(A_j\ell - A_jA_j^*w)$ and their feature

difference is $\sum_j (B/r_j)(\|\ell\|^2 - w^\top A_j\ell)$. Let $\bar{D}$ bound the mean-distance norm. Multiplying both densities by $\lambda = \min(1, \tau/\bar{D})$ (one if $\bar{D} = 0$) makes them feasible for Equation (45). This scaling preserves the original class when $\|\rho_0^{\text{pilot}}\| \leq B$, because zero is in the ball and the ball is convex. Thus the scaled feature gap supplies a constructive lower bound even if optimization is incomplete.

We minimize Equation (46) by iteratively reweighted quadratic solves and retain both the best feasible pair and the best dual weights, which may occur at different iterations. Continuous Gram integration errors pad each residual norm and the mean-distance norm. Translation is a real orthogonal rotation of each Fourier cosine/sine pair, so the phase-wrapped Gram and target preserve those error norms. The witnesses are continuous target/adjoint fields plus the known pilot; they are not restricted to a voxel representation. Floating-point magnitude guards are additional numerical checks, not validated interval arithmetic.

The prespecified development grid uses the existing 128-particle, radius-5 geometries and broad targets on three stacks. It includes identical nominal poses, antipodal coherent rotations, and antipodal random joint-ball poses at one and two degrees with $0.5 \text{ \AA}$ translation radius. The same random unit vectors are used at both angular budgets. Independent conic programs check small Hilbert-space cases; deliberately crude quadrature is checked against the exact continuous sinc Gram; nonlinear constant-cell projections check the phase convention. A serialization failure and the first complete grid remain archived. The final rerun adds the numerical guards and separate witness records requested by the focused mathematical audit.

The fixed-pair problems are solved to relative brackets below 0.5%. Their lower bounds remain substantially smaller than several pose-aware upper intervals. This leaves a gap: the selected configurations need not maximize ambiguity, so these results do not prove that the larger upper bounds can be attained or substantially reduced. The testing statement also retains the declared white-Gaussian noise, density and nuisance assumptions.

A bounded alternating ambiguity probe. We next test whether changing the selected poses strengthens a witness, using one prespecified 10049 central-feature case at two degrees and $0.5 \text{ \AA}$. Project the two unscaled continuous densities onto a nested constant-cell basis containing the original pilot. Orthogonal projection cannot increase their distance from that pilot, so they remain in the original continuous class. We hold the projected densities fixed and minimize their whitened mean-image distance over both allowed pose collections. Analytic cell Fourier moments provide derivatives, while independent direct cell sums check every particle at the saved feasible poses. A subsequent

Stack / fit	Target	Degrees	Joint/sharp	Pilot	Stress / upper	Min. power
10049	center	0.5	0.194	0.175	0.803	0.100
10049	center	1	0.291	0.255	0.630	4.4e-09
10028	center	0.5	0.124	0.110	0.864	1.000
10028	center	1	0.189	0.162	0.638	0.005
10028	center	2	0.495	0.436	0.274	0 [†]
10049	center	2	0.647	0.568	0.313	9.9e-114
10049	contrast	0.5	0.225	0.206	0.828	3.4e-26
10049	contrast	1	0.327	0.293	0.668	1.2e-63
10049	contrast	2	0.715	0.633	0.335	0 [†]
10076	center	0.5	0.138	0.127	0.865	0.997
10076	center	1	0.202	0.179	0.637	0.214
10076	center	2	0.472	0.427	0.360	3.3e-72
10076	contrast	0.5	0.166	0.152	0.866	6.9e-20
10076	contrast	1	0.245	0.214	0.650	7.2e-41
10076	contrast	2	0.578	0.512	0.323	8.5e-238
10049 / full	center	2	0.557	0.510	0.667	4.9e-122
10049 / cuts	center	2	0.572	0.512	0.636	4.0e-109
10049 / sharp fit	center	2	0.548	0.503	0.630	9.3e-119

Table 4. Known-pilot refinement on the same weights and spectral event as the joint/sharp bound. Both retain the full continuous density class, a 0.07-field Gaussian width and the joint pose ball with 0.5 Å translation radius. Widths are relative to the no-data bound; stress and power use the same conditional reference scenarios as Table 3. A width improvement need not yield feature detection (†: probability underflow).

Stack	Target	Nominal	Coherent 1°	Random 1°	Coherent 2°	Random 2°
10028	center	0.087	0.087	0.088	0.087	0.088
10028	contrast	0.106	0.108	0.107	0.113	0.107
10049	center	0.123	0.123	0.125	0.123	0.126
10049	contrast	0.147	0.148	0.148	0.151	0.150
10076	center	0.106	0.106	0.106	0.106	0.106
10076	contrast	0.126	0.127	0.126	0.132	0.127

Table 5. Constructive lower bounds on the half-width of any uniformly valid deterministic-length interval, divided by the no-data half-width $B \|\ell\|$. Each entry uses two feasible antipodal pose configurations, continuous densities in the original ball and a Gaussian testing distance below $2\Phi^{-1}(0.95)$. Random poses are fixed design draws. All 30 fixed-pair modulus gaps are below 0.5%; this does not certify maximization over poses. A small entry cannot establish that a larger upper bound is loose.

density step returns to the full continuous modulus problem at those poses. Each witness is feasible without a global pose-optimization claim.

The initial 24-, 48- and 72-cell pose steps give lower widths 0.09624, 0.12169 and 0.12438 of no data, all below the source continuous witness’s 0.12634. Cell projection changes the images and loses target energy; these outcomes are retained. Once the mean distance falls below the testing budget, pose fitting alone cannot enlarge that fixed pair’s feature gap. Refitting the continuous densities at the 72-cell step’s poses raises the lower width to 0.13884. The two further alternations allowed by the development protocol raise it to 0.14594 and 0.14962, with final fixed-pair bracket gap 0.1422%. We stop at the specified three rounds. This is a stronger selected witness, not a global ambiguity limit. The

projected densities carry substantial negative-value energy (approximately 0.49 and 0.36 of total squared norm in the 72-cell pair), illustrating why the stated signed class must not be confused with biological admissibility.

E.4. An experimental noise-calibration attempt

The following elementary conservative calculation allows nonzero calibration signals. Suppose independent $Y_i \sim N(\mu_i, \Sigma)$ have arbitrary fixed means and a common covariance. Put $T = \text{tr } \Sigma$ and $S = \sum_i \|Y_i\|^2$. For $\lambda > 0$,

$$\begin{aligned} \mathbb{E} e^{-\lambda S} &\leq \det(I + 2\lambda \Sigma)^{-n/2} \\ &\leq (1 + 2\lambda T)^{-n/2}. \end{aligned}$$




Figure 3. A continuous testing witness after the first alternating pose/density refit on 10049, displayed as 48-cell averages of the two central slices. The common color scale is symmetric and clipped. These are constructed signed densities in the declared mathematical class, not experimental reconstructions or a pair asserted to be molecular structures. Their allowed poses and whitened image-mean distance yield the conditional fixed-length lower bound; the displayed projections are for visualization only.

The first inequality follows from the noncentral Gaussian Laplace transform; the second uses $\prod_j (1+x_j) \geq 1 + \sum_j x_j$ for nonnegative eigenvalues. Optimizing the exponential Markov inequality gives

$$\Pr\{S \leq ntT\} \leq \exp\{n(1-t+\log t)/2\}, \quad 0 < t < 1.$$

Choose t to make the right side β . Then $T_+ = S/(nt)$ is an upper bound on T with probability at least $1 - \beta$, and $\Sigma \preceq T_+ I$ on that event. This can be much more conservative than estimating a known diagonal noise model. Independent inference with the same covariance permits allocating $\alpha_{\text{noise}} + \beta \leq \alpha$. The result does not prove the covariance commonality or pose-independence assumptions for archive data.

The exploratory three-stack attempt fits each old pilot’s scalar amplitude on pilot-pool particles, then calibrates from one particle per new exposure group. Half of the new groups calibrate, and the others supply energy diagnostics. Inference also uses one particle per old inference exposure group. No background mask or fitted-signal subtraction enters the calibration statistic. We use $\beta = 0.005$, $\alpha_{\text{noise}} = 0.045$, density radius two, and prescribed 0/1/2-degree sensitivities with 0.5 Å shifts at nonzero angles.

All eighteen intervals contain their amplitude-aligned deposited-map value; none excludes zero. This is a precision failure, not evidence of density calibration. The 10 Å targets revert to no data at nonzero pose budgets. The earlier 128-particle attempt on 10028 is also retained; it additionally requires unverified independence within repeated exposure groups.

Independent calibration in the estimator’s directions.

The trace envelope protects every direction, although a fixed feature estimator needs only one scalar variance. Let V

contain its raw-coordinate particle weight rows $v_i^\top$, chosen independently of a second calibration pool. Under the same common-covariance and independence assumptions, its noise variance is $s^2 = \sum_i v_i^\top \Sigma v_i = \text{tr}(V \Sigma V^\top)$. Applying the preceding lemma to VY_j gives the following direct corollary:

$$s^2 \leq s_+^2 := \frac{\sum_j \|VY_j\|^2}{nt} \quad \text{with probability at least } 1 - \beta.$$

The projected coordinates may be correlated. Combining this event with the bias-aware critical width at α_{noise} gives pointwise failure at most $\beta + \alpha_{\text{noise}}$. Indeed, the width is at least the bias bound, and its Gaussian outside probability increases with noise scale. This is an elementary specialization of the trace argument, not a new concentration principle or a simultaneous guarantee for selected targets.

For the same calibration pool, $\sum_j \|VY_j\|^2 \leq \|V\|_F^2 \sum_j \|Y_j\|^2$, so directional calibration cannot exceed the original trace envelope for these weights. To preserve independence here, we retain weights selected using the original first exposure half and calibrate their variance with the second half. The pools therefore differ in the empirical comparison. Before projection, weights are pulled back through each particle’s known centering phase: complex coefficient $w_R + iw_I$ becomes $\bar{p}(w_R + iw_I)/\sigma_{\text{old}}$ in raw Fourier coordinates. Ignoring this step would be incorrect for colored noise.

The new standard-deviation bounds are 0.132–0.285 times the original bounds. One broad fixed-pose interval excludes zero; none does at one or two degrees. Thus reducing this avoidable noise conservatism does not resolve the pose/density precision problem. No experimental density-coverage inference follows from deposited-map agreement or this conditional fixed-pose detection.

Stack	$n_{\text{inf}}/n_{\text{cal}}$	σ_ℓ (Å)	Center $\pm$ half-width	Ref.
10028	58/57	33.8	0.94 ± 5.52	0.63
10049	30/35	16.5	3.06 ± 14.23	1.02
10076	93/92	29.3	4.37 ± 7.49	2.08

Table 6. Exploratory experimental sensitivity intervals for the broad central Gaussian average at 1° and 0.5 Å translation radius. Density units use the separately amplitude-aligned unit-norm pilot. One particle per inference/calibration exposure group is used; remaining new groups supply diagnostics. None excludes zero. Ref. is a pilot-pool amplitude-aligned deposited-map value, not a coverage label. Noise commonality, consensus-pose independence and density/pose radii remain unverified.




Figure 4. Experimental sensitivity intervals with one inference particle per exposure group. Kernel widths are Gaussian standard deviations, not FSC resolutions. None of these intervals excludes zero. All 10 Å targets switch to the no-data center and width at nonzero pose budgets. Supplied density/pose classes and unverified common-noise/consensus-pose assumptions preclude an experimental coverage claim.

Refitting with an independent covariance guide. We also use the first exposure half’s uncentered second moment $\hat{C}$ to form $C = 0.8\hat{C} + 0.2\text{tr}(\hat{C})I/m$; the second half still supplies the final directional noise bound. This positive matrix guides optimization only. With known centering matrices O_i , set $W = \text{diag}_i[(O_i C O_i^\top)^{-1/2}]$ and write $w = Wx$. The continuous fixed-pose problem then uses target vector $W\mathcal{A}\ell$ and Gram WGW . Its proxy objective is $z\|x\| + B\|\ell - \mathcal{A}^*Wx\|$. Quadrature residual errors are evaluated at Wx ; transformed Gram-action errors multiply by $\|W\|$. None of these proxy quantities replaces the independent final noise or pose audit.

All six proxy fits reach a sum-objective gap below 0.33%. Their fixed-pose intervals improve, but five one-degree intervals widen and the sixth remains vacuous. Nine of eighteen settings now use no data, including every two-degree setting. Only the same broad 10076 fixed-pose interval excludes zero. Thus better noise-guided fixed-pose fitting can amplify pose sensitivity; it does not substitute for joint pose-aware design.

A larger physical 10 Å target. On 1,024 particles from 10049 at Fourier radius 12, with target Gaussian standard deviation 10 Å and sampled band ending at 19.68 Å, an

order-80 continuous fixed-pose fit reaches sum-objective gap 0.000992 in 3,583 seconds, with relative half-width 0.01378. The independent normalized reference-map central value is 1.9245; conditional sign power is one at fixed pose. Keeping those weights and auditing a joint ball of one degree and 0.5 Å changes the conclusion: even with sharp cube constants, residual cross terms and the known-pilot refinement, the result uses the no-data interval and has zero sign power. The polynomial pose term is 6.376 and the cubic remainder is 56.493; the final unselected half-width is 61.223, exceeding the no-data half-width 34.390. The pose audit takes 1,491 seconds and 1.08 GB peak resident memory. Thus increasing the particle count and band does not by itself control the nonlinear remainder. These are conditional development results under supplied noise and nuisance bounds; the 128- and 1,024-particle subsets are not nested, and the latter fit is not pose-optimized.

F.5. Fourier cancellation in the nonlinear remainder

The higher-band failure motivates a further exploratory refinement of the *same* continuous-density class. The preceding frequency-wise Taylor bound discards cancellation in the known adjoint field. For one particle let $f(x) = \text{Re} \sum_q c_q \exp(2\pi i k_q^T x)$, where $c_q = (w_{Rq} + iw_{Iq})C_q/\sigma$.

Stack	New cal. n	SD ratio	Center $\pm$ half-width	Ref.
10028	58	0.250	0.94 ± 4.62	0.63
10049	36	0.180	3.06 ± 8.94	1.02
10076	93	0.285	4.37 ± 5.08	2.08

Table 7. Independent directional-noise re-audit of the same broad experimental estimators at 1° and $0.5 \text{ \AA}$. SD ratio compares the new noise bound with the original first-pool trace bound; the new second-pool sample differs. Centers and bias bounds are unchanged before no-data selection. None excludes zero at this pose budget. Ref. remains an approximate deposited-map diagnostic.




Figure 5. All eighteen experimental directional-noise sensitivity cases. Only the calibration bound changes; weights and continuous bias terms are held fixed. The single interval excluding zero is the broad 10076 target at fixed pose; it does not exclude zero at one degree. Noise, density and pose assumptions remain unverified, and the previously inspected diagnostic pool makes this exploratory reuse rather than a new frozen confirmation.

Let $R = \sqrt{3}/2$, let a be the rotation radius, and let s_i bound the norm of the three-dimensional translation dual to the detector shift. For exact isometric detector embeddings, $s_i = s$; in computation we use the embedding pseudoinverse and an explicit residual-phase correction, as recorded with the implementation.

For a joint scaled pose ball, the path $y(t) = \exp(t\Omega)x + tb$, $0 \leq t \leq 1$, satisfies

$$L = \sqrt{a^2 R^2 + s_i^2}, \quad H = a^2 R, \quad T = a^3 R \quad (48)$$

as bounds on its first, second and third derivative norms. An enclosing ball has radius $R + s_i$. An alternative enclosing coordinate cube has half-side

$$r_i = \frac{1}{2} + 2R \sin\{\min(a, \pi)/2\} + s_i. \quad (49)$$

Indeed, $\|(\exp(t\Omega) - I)x\| \leq 2R \sin\{\min(a, \pi)/2\}$, and translation adds at most s_i . This enlarges the integration

domain of a *known* field; it imposes no new support or smoothness condition on the unknown density.

Proposition 4 (Enclosing-domain remainder). *Let D_i contain every transformed unit cube along every allowed path, and let $S_m = \|D^m f\|_{L^2(D_i), F}$ be the Frobenius derivative-tensor norm. The second-order pose Taylor remainder has L^2 norm at most*

$$r_{\text{enc},i} = \frac{1}{6} \{L^3 S_3 + 3LH S_2 + T S_1\}. \quad (50)$$

Proof. The third chain-rule derivative is the sum of the cubic derivative contraction, three Hessian contractions with y' and y'' , and the gradient contraction with y''' . Cauchy–Schwarz bounds each contraction by its Frobenius norm. The map from x to $y(t)$ is rigid and has Jacobian one; enlarging its image to D_i increases each nonnegative squared integral. Minkowski’s inequality and Taylor’s integral remainder give (50). $\square$

Stack	σ_ℓ (Å)	Fixed / old	Fixed / new	1° / old	1° / new
10028	33.8	0.111	0.098	0.286	0.751
10028	10.0	0.907	0.899	1.000	1.000
10049	16.5	0.398	0.272	0.553	0.905
10049	10.0	0.728	0.625	0.853	1.000
10076	29.3	0.158	0.133	0.314	0.566
10076	10.0	0.894	0.867	0.971	1.000

Table 8. Relative half-widths before and after covariance-guided fixed-pose refitting. Both variants independently calibrate noise on the same second exposure half. A training second-moment proxy guides the new weights; it is not a noise confidence bound. All six fixed-pose widths improve, but five one-degree widths worsen and the sixth remains vacuous. At two degrees every new fit falls back to no data.

For either symmetric domain, write $K_D(v) = \int_D e^{2\pi i v^T x} dx$. The exact real-arithmetic identity retaining cross-frequency terms is

$$S_m^2 = \frac{1}{2}(2\pi)^{2m} \operatorname{Re} \sum_{q,l} (k_q^T k_l)^m [c_q \bar{c}_l K_D(k_q - k_l) + (-1)^m c_q c_l K_D(k_q + k_l)]. \quad (51)$$

The sum over ordered tensor derivatives yields $(k_q^T k_l)^m$; the sum-frequency sign is $i^{2m} = (-1)^m$. For a ball of radius r , $K_D(v) = 4\pi r^3 j_1(z)/z$, $z = 2\pi r \|v\|$, with value $4\pi r^3/3$ at zero. For the coordinate cube, $K_D(v) = \prod_j 2r \operatorname{sinc}(2rv_j)$. The scalar density bias receives $(B + P) \sum_i r_{\text{enc},i}$. Taking the smaller valid remainder for each particle is deterministic and uses the original spectral event; no new random certificate is selected. Magnitude-based roundoff guards remain heuristic, not validated arithmetic.

Independent spherical and tensor-product quadrature checks all ordered derivative tensors through degree three. Direct nonlinear cube integration also checks allowed joint poses and deliberately nonplanar frequency embeddings. On the saved 1,024-particle higher-band example, the ball enclosure reduces the cubic contribution from 56.493 to 32.623 after particle-wise selection, still yielding no data. The subsequent expanded-cube probe reduces it to 21.110, with total half-width 25.840, or 0.751 of no data. All particles improve in that probe; its extra embedding-residual bias is below 5×10^{-14} . Both probes take approximately five CPU seconds and retain the previous polynomial certificate and estimator weights. Nevertheless, the three reference scenarios still have numerically zero feature-detection power. This inexpensive refinement reduces a conservative remainder; it does not solve the fine-feature usefulness or experimental calibration problem.

E.6. Uniform CTF sensitivities

For fixed voltage, spherical aberration and amplitude contrast a , $C = \sin(\gamma - \arcsin a)$. If each defocus axis can change by d and its astigmatism angle by b radians, a conservative defocus error is $d + |\text{df}_U - \text{df}_V| \sin(\min(b, \pi/2))$.

Multiplying by $\pi \lambda s^2$ and adding the phase-shift radius bounds the phase error. Evaluate the exact sine range on that interval, then multiply by the positive relative gain/B-factor range. This yields an absolute transfer envelope e_{iq} without division by a CTF zero.

For any allowed pose, a complex Fourier wave has unit modulus on the unit-volume cube. Therefore the additive CTF bias is bounded by

$$b_{\text{CTF}} \leq (B + P) \sum_{iq} |w_{iq,\mathbb{C}}| e_{iq} / \sigma. \quad (52)$$

This term can be combined with the nominal-CTF pose bound, including simultaneous errors, but discards cancellations and is not optimized jointly with them. Table 9 applies it to all six saved radius-12 fixed-pose fine-feature fits. Defocus uncertainty of 100 Å already causes two no-data fallbacks; all six fall back at 500 Å and under the combined sensitivities. These negative results diagnose a loose audit, not an impossibility theorem for other estimators or a calibrated distribution of experimental CTF error.

E.7. An explicit support ablation

Changing support changes the density class. For a centered cube of side $s \leq 1$, the transformation $\tilde{\rho}(u) = s^{3/2} \rho(su)$ is an L^2 isometry. The transformed acquisition uses $k' = sk$, $q' = sq$, $C' = s^{3/2}C$ and translation radius $t' = t/s$. A normalized Gaussian target uses center c/s , width σ_ℓ/s and amplitude $s^{-3/2}$. Thus the target and nonlinear acquisition values are preserved, including physical translations.

Hard support need not be imposed without a sensitivity allowance. Write $\rho = p_S + d_S + r_{S^c}$, where p_S, d_S are supported inside the smaller cube, $\|d_S\| \leq B$ and $\|r_{S^c}\| \leq \epsilon$. Add to the inside audit

$$b_{\text{tail}} \leq \epsilon \left[\|A_0^* w - \ell\| + \sum_{iq} |c_{iq}| d_{iq} \right], \quad d_{iq} = \min \left(2, 2\pi \sqrt{\theta^2 \|k_{iq}\|^2 / 12 + t^2 \|q_{iq}\|^2} \right). \quad (53)$$

CTF sensitivity	Relative half-width	No-data cases
Nominal	0.120–0.238	0/6
Defocus ± 100 Å	0.647–1.000	2/6
Defocus ± 500 Å	1.000–1.000	6/6
Combined nuisance bounds	1.000–1.000	6/6

Table 9. Declared CTF-error envelopes added to all six fixed-pose radius-12 fine-feature fits. The combined setting is 100 Å defocus, 2° astigmatism angle, 1° phase, 5% gain and 10 Å² relative B-factor. Radii are sensitivities, not calibrated errors. The triangle bound discards cancellation; its vacuity is not an impossibility result for other methods.

Indeed, the spatial phase has zero mean on the cube, its integrated squared linear part is $\|\Delta k\|^2/12$, and the rotation chord is at most $\theta \|k\|$. Apply $|e^{ia} - e^{ib}| \leq \min(2, |a - b|)$ and the triangle inequality, then restrict the residual to S^c . The bound is conservative; it does not infer ϵ from the data.

The initial radius-12, 10 Å sigma support probe on 10049 uses 128 particles, 1° and 0.5 Å translation. Side $s = 0.5$ gives relative width 0.417 under hard support, yet the cropped-reference minimum sign power is below 10^{-150} . A tail radius of 0.1 causes no-data fallback. The original unit-norm reference has outside norm 0.320 and is excluded by all these tail settings; it is not relabeled as an in-class validation example. The completed full-cube control takes 6,324 seconds, attains a fixed-pose sum-objective gap of 0.19%, and falls back to no data at every tested tail allowance after adding pose sensitivity. Its integrated cubic term alone is 115.4 in the normalized density units. A converged nominal fit therefore does not establish a useful pose-aware interval.

F.8. A closer Fourier variational comparison

We independently implement the fixed-pose Gaussian variational objective of Ullrich et al. (2020), Equation 14, in a Hermitian Fourier grid with trilinear slice interpolation. The original authors’ full software and pose-learning pipeline are not reproduced. CTFs, an independent pilot as prior mean, physical target normalization and a declared prior-scale sweep are explicit extensions. For $x \sim N(m_0, \tau^2 I)$ and whitened $y = Ax + e$, put $Q = A^\top A + \tau^{-2} I$. The diagonal-Gaussian optimum has the full posterior mean and coordinate variances $1/Q_{jj}$. A target’s full posterior variance is $\ell^\top Q^{-1} \ell$, whereas its diagonal approximation is $\sum_j \ell_j^2 / Q_{jj}$. Differentiating the Gaussian objective verifies these solutions without stochastic-training error. The identity

$$\|w\|^2 + \tau^2 \|A^\top w - \ell\|^2 = \ell^\top Q^{-1} \ell, \quad w = A Q^{-1} \ell,$$

separately checks prior-predictive calibration.

Each stack uses 1,024 particles, radius 12 and two 10 Å-sigma targets: a central average and a contrast between centers at $z = \pm 10$ Å. The prior expected squared deviation norms are 0.25, 1 and 4; all settings are retained. The sparse conditional calculations take 3.3–3.9 seconds per

stack and 45.0 MB per acquisition matrix. The full posterior’s matched prior-predictive coverage is 0.95, and the variance identity error is below 4×10^{-16} . Diagonal feature intervals are *wider* in every tested case, by factors 1.073–1.415, with matched prior-predictive coverage 0.964–0.994. Diagonal covariance is therefore not a general explanation for undercoverage of these targets.

Table 10 and Figure 6 instead report fixed-map, known-noise coverage. Interpolation and finite-representation forward discrepancy is 16.8%, 16.2% and 25.3% across the three reference generators. Together with prior shrinkage, this is a material baseline limitation. These controls do not refute Bayesian prior-predictive guarantees or constitute experimental density-coverage measurements. All matched-model and nonlinear pose controls are retained in the public records.

We also audit each of these 18 estimators in the continuous cube without changing its original affine center. Let κ be the known difference between its interpolated-prior center and the exact continuous-pilot center. The fixed-pose bias bound is then $|\kappa| + 2 \|\ell - A^* w\|$. Audited half-widths are 0.257–0.815 of no data and 56–352 times the original full posterior widths. The nominal continuous-map checks are covered conservatively, but sign power is numerically zero throughout, even after correcting κ . Thus the broad continuous class already loses useful precision before adding pose uncertainty. This does not establish impossibility for a physically narrower class or a different estimator.

G. Uncertainty Targets in the Recent Literature

The following distinctions guide baseline selection. They do not rank methods that solve different problems on one inappropriate score. The accompanying dated survey has 101 curated candidates across reconstruction, poses, heterogeneity, map validation, atomic ensembles and general inverse inference; its source manifest distinguishes retrieved material from targeted critical reading and inaccessible full text. This is a scoped critical survey rather than a claim to have exhaustively reviewed all cryo-EM applications. Versioned preprints are identified separately from published articles.

Continuous Fourier-Slice Audits for Cryo-EM Density Uncertainty

Stack	Target	VI/full width	Full coverage	VI coverage
10028	Center	1.265	0.032	0.162
10028	Contrast	1.154	0.000	0.000
10049	Center	1.176	0.000	0.000
10049	Contrast	1.245	0.000	0.000
10076	Center	1.205	0.000	0.000
10076	Contrast	1.208	0.728	0.898

Table 10. Fixed-pose Fourier Gaussian baselines at prior expected squared deviation norm one, 1,024 particles per stack, radius 12 and 10 Å target sigma. Coverage integrates known measurement noise for the continuous deposited-map generator at nominal poses; it is not empirical experimental coverage. The full posterior has exactly 0.95 coverage under its matched prior predictive distribution. VI is wider for these targets. Poor fixed-map coverage includes shrinkage and interpolation/representation error; it does not refute the source method's stated probabilistic assumptions.

Additional local and atomic validation. The recent Gold-standard local toolbox compares half-map phase and amplitude, with explicit independent-noise assumptions in its SNR derivation (Vargas et al., 2026). EMMIVox combines molecular force fields, map likelihoods and half-map-informed uncertainty for atomic refinement (Hoff et al., 2024). Neither output is the fixed raw-particle density functional studied here.

Direct statistical uncertainty prior art. The probabilistic Fourier-slice model of Ullrich et al. (2020) is a close reconstruction-uncertainty comparator and explicitly recognizes model bias. Joint likelihood soft modes (Rangan et al., 2024) expose pose-density ambiguities. Lai et al. (2023) connect empirical Bayes and sequential MCMC filtering to dynamic image reconstruction, including scalar confidence sets and hybrid resampling. Their approximate fitted-family resampling and posterior computations differ from the present uniform conditional bounded-class statement; they cannot be dismissed as deterministic reconstruction methods. Neither general bias-aware inference nor the distinction between posterior variance and model error originates here.

Adjacent radiative CT work also separates analytic Gaussian uncertainty from shared reconstruction error and distinguishes foreground evaluation from whole-volume scores (Zhao et al., 2026). Its density-moment formulas and validation distinctions further constrain novelty claims; its acquisition and estimands differ from cryo-EM.

Pose inference versus pose confidence. Posterior-averaged orientation estimation (Xu et al., 2026) and probabilistic orientation search in SIMPLE (Van et al., 2025) improve pose estimation through different statistical or optimization constructions. Neither an orientation distribution nor a stochastic search neighborhood should be relabeled a calibrated per-particle confidence ball without verifying its guarantee. Our prescribed local balls avoid that claim, but consequently do not solve practical pose calibration.

Cryo-forum uses quaternion-distribution dispersion for er-

ror ranking and filtering (Chung, 2023). ARCHER learns reference-conditioned pose scores and evaluates cross-specimen transfer and junk rejection (Nguyen & Pham, 2026). CESPED standardizes experimental poses obtained by refinement, and explicitly acknowledges uncertainty in those labels (Sanchez-Garcia et al., 2024). CryoPARES transfers alignments between related experiments (Sanchez-Garcia et al., 2026); CryoFastAR combines multiview pose prediction, synthetic training and experimental fine-tuning (Zhang et al., 2025b). These approaches address efficiency and different transfer settings. Their quality scores and angular errors cannot supply a uniform maximum-error radius for every particle in a new stack. The 0.1° and 0.5° budgets in our frozen study are sensitivity settings, not experimentally established alignment accuracy.

Physical heterogeneity versus estimation error. RECOVAR estimates covariance and uses kernel regression to account for noisy latent coordinates (Gilles & Singer, 2025). DynaMight and 3DFlex represent motion (Schwab et al., 2024; Punjani & Fleet, 2023). Particle-counting work explicitly studies when assignments misestimate state populations (Evans et al., 2026). These motivate separate tests of population recovery, structural variation and uncertainty in a homogeneous density. The present bounded-density target covers only the last. CryoDRGN-AI addresses joint ab initio reconstruction (Levy et al., 2025); fixed-pose homogeneous training does not reproduce that task.

DynaMight's held-out deformation disagreement and independent-half training assess motion errors; 3DFlex separates lower-frequency motion training from higher-frequency reconstruction validation. Their protocols must be examined at the level of shared components and evaluated frequencies. Mattingly et al. (2026) study how the image model limits learnable conformational populations, using information-based selection of ensemble granularity. That prior-averaged population target differs from our fixed homogeneous density functional. CryoBench likewise distinguishes representation learning, conformation assignment and volume recovery with several reference-based FSC met-




Figure 6. Known-noise coverage for fixed continuous reference densities at nominal poses as the declared Gaussian prior scale varies. The dashed line is 0.95. This is distinct from the full posterior’s matched prior-predictive coverage, which is 0.95 throughout. Diagonal VI is wider here; prior and forward-model discrepancy remain important. No prior scale is selected using these outcomes.

rics (Jeon et al., 2025). Our simulations posit one density per acquisition geometry; the actual stacks need not be homogeneous. In particular, the assembly-intermediate mixture in EMPIAR-10076 does not validate our shared-density assumption.

Map validation and learned priors. False-discovery confidence maps test density significance against a noise reference (Beckers et al., 2019). Directional resolution measures acquisition anisotropy (Vilas et al., 2020); independent-particle tests assess map fit (Ortiz et al., 2020). Bias/variance analyses explain why validation requires more than internal agreement (Sorzano et al., 2022). Blush introduces learned reconstruction regularization (Kimanian et al., 2024), while LocScale-2.0 attaches confidence to enhancement relative to its stated proxy (Bharadwaj et al., 2026). Compressed-image recovery with cryoSENSE targets acquisition fidelity rather than density-confidence coverage (Shabeeb et al., 2026). Matched evaluations must retain these targets. A sharper-looking map or a significant positive voxel is not itself a bound on error relative to unknown density.

CryoRes predicts local resolution using deposited and algorithm-derived training labels (Dai et al., 2023). Such supervision inherits the meaning of those labels; it does

not transform a resolution score into a confidence interval. Similarly, held-out image prediction measures agreement with new observations, and can remain favorable when two reconstructions share unsupported density directions. Ground-truth simulation, new-exposure prediction, directional diagnostics and independent structural evidence therefore complement one another.

Invariant compression and thermodynamic targets. Moment-based posterior sampling conditions a learned prior on shift-invariant power spectra in one-dimensional orbit recovery (Janson & Andén, 2025). Compression discards phase information, so posterior concentration can rely on the prior. Separately, cryoTWIN/PaStEL targets conformational free energies under sampling, pose and equilibrium assumptions (Yamazaki et al., 2026). Neither problem is equivalent to our shared-density feature interval.

Atomic ensembles and independent modalities. Experiment-guided AlphaFold3 conditions structural generation on reconstructed ESP maps and other measurements (Maddipati et al., 2026). Its atomic-ensemble target differs from raw-particle density inference. The multimodal examples also illustrate why improved map agreement need not imply improved fit to a separate measurement modality.

Independent evidence is relevant to assessing structural interpretations that the density measurements alone leave ambiguous.

Coordinate-error conformal prediction. CalPro calibrates protein-coordinate error scores at an exchangeable-protein level (Shihab et al., 2026). Its atomic target and calibration unit differ from a fixed density feature inferred from noisy projections. We do not interpret a calibrated protein-average score as simultaneous residue coverage or as experimental-density coverage.

Simultaneous inverse inference. Constrained test-inversion methods already address joint confidence regions for multiple functionals (Batlle et al., 2025). A single data-consistency region can support simultaneous projections, whereas separately calibrated marginal intervals generally require an appropriate correction. Our present scalar-target theorem is not a new general framework for simultaneous inference, and a colored map of marginal intervals must not be presented as simultaneous coverage.

Conformal calibration and its sampling unit. Posterior-variance calibration can provide error bounds even when a posterior sampler is approximate (Narnhofer et al., 2024); its theorem uses independent identically distributed signal-observation pairs and conditions on a variance bin. QUTCC learns quantile-dependent image predictions and calibrates marginal pixel errors from paired reference images (Ye et al., 2026). These methods can be useful without our fixed-signal radius assumption, but require a calibration population with the stated sampling relationship to deployment. Voxels or noise draws from one molecule are not automatically independent calibration molecules. Neither marginal population coverage nor variance-bin coverage is the same as coverage uniformly over every density in a specified class.

Prior-assisted alignment and atomic refinement. CryoPROS uses generated auxiliary particles for pose correction, with template-variation and noise-only controls (Zhang et al., 2025a). Such tests address model bias without calibrating experimental pose confidence regions. CoCoFold fine-tunes AlphaFold against particle images using a Gaussian-mixture forward model and supplied poses/CTFs (Liao et al., 2026). Its pose-perturbation checks support a sensitivity comparison, not a density-coverage guarantee. Both motivate independent validation of prior-assisted inference; neither is a matched density-functional uncertainty baseline here.

Open questions suggested by the survey. Particularly consequential gaps are absolute pose-confidence calibration under reference error; uncertainty in conformational populations when candidate structures are imperfect; validation

Table 11. Experimental-image prediction error on development test groups (normalized MSE; lower is better). Neural checkpoints minimize tuning-group error among completed runs: epochs 60, 50 and 100 respectively. All half-map FSCs remain above 0.143 through the sampled limit in the last column, which is not a measured crossing.

EMPIAR	Gaussian	Voxel	Neural	Limit (Å)
10028	0.8355	0.8417	0.8293	16.08
10049	0.9308	0.9316	0.9302	7.87
10076	0.8764	0.8806	0.8707	13.97

of detail introduced by learned priors using independent particles; and uncertainty propagation through CTF fitting, image preprocessing, alignment and atomic interpretation. Acquisition chosen for a specific structural feature is another possible application of adjoint sensitivity. These are directions inferred from the literature, not priority claims or completed capabilities of this implementation. The present continuous audit addresses representation exclusion under stated nuisance assumptions; solving those assumptions empirically remains distinct work.

H. Additional Development Comparisons

H.1. Experimental reconstruction and prediction context

We execute stock cryoDRGN 4.3.1 homogeneous neural training on both exposure-group halves of all three stacks: three hidden layers, width 256, batch size eight, fixed consensus poses and CTFs. These are actual neural fits; they are not ab initio inference or uncertainty estimates. Gaussian and voxel ridge fits use the same halves. All methods are evaluated on identical held-out particles and 1,410 Fourier pairs per image inside radius 30 of the 64-pixel grid, with the same window and image-background normalization. Neural training itself uses the stock radius-32 loss and per-half global scaling, so training objectives are not identical.

Table 11 selects checkpoints using tuning-group error only, from runs through 60 epochs for EMPIAR-10028/10049 and 100 epochs for EMPIAR-10076. The selected epochs are 60, 50 and 100. Two are the last available checkpoints, so undertraining remains possible. The development plateau check motivated these budgets; it is not a preregistered stopping rule or a proof of global convergence. Exposure-group bootstrap comparisons, checkpoint hashes and direct-network versus map-Fourier checks are recorded with the results. FSC assesses conditional map agreement; neither it nor image prediction labels true experimental density coverage.

A frozen follow-up evaluates these unchanged models on 4,096 additional particles per dataset from exposure groups

Table 12. Joint ambient/pose interval width divided by the no-data width. Columns 2–4 use a 100-iteration cap. The last column repeats 2° with up to 500 iterations; EMPIAR-10049 retains a 0.548% gap, above the 0.5% tolerance. Others meet tolerance. These are different optimization budgets, not separate independent replications.

EMPIAR	0.1°	0.5°	2°	2° , longer
10028	0.087	0.200	0.866	0.865
10049	0.118	0.239	0.996	0.992
10076	0.110	0.218	0.910	0.907

absent from development (Table 13). Neural prediction has lower NMSE than both controls in all three datasets; all six prespecified paired group-bootstrap differences remain positive under the stated multiplicity adjustment. This confirms conditional prediction ranking on new exposures, not experimental density coverage or independent-pose FSC.

Post-hoc background diagnostics on all 12,288 new images show mean nearest-neighbor normalized pair products of 0.197–0.292 across the three stacks and two image axes (Figure 10). White-Gaussian controls with the same background centering and normalization are near zero. The annular background can contain molecular signal, ice and preprocessing artifacts, so this is not identification of the measurement-noise law. It does show why simple spatial scaling cannot be assumed to establish whitening; our supplied-noise simulations do not resolve that experimental issue.

Table 13. Frozen additional-exposure prediction: 4,096 particles per dataset from 115, 71 and 185 source groups absent from development. Models and cohort were locked before image access. Lower normalized prediction MSE is better. Supplied consensus poses/CTFs remain conditional; this does not measure unknown density coverage.

EMPIAR	Gaussian	Voxel	Neural
10028	0.835319	0.841379	0.829369
10049	0.931579	0.932359	0.930739
10076	0.874799	0.879078	0.868657

EMPIAR	Comparator minus neural	NMSE difference	Adjusted percentile interval
10028	Gaussian	0.005950	[0.004995, 0.006932]
10028	Voxel	0.012011	[0.010695, 0.013356]
10049	Gaussian	0.000840	[0.000659, 0.001023]
10049	Voxel	0.001620	[0.001366, 0.001875]
10076	Gaussian	0.006143	[0.005686, 0.006593]
10076	Voxel	0.010421	[0.009792, 0.011038]

All six prespecified paired differences are shown. Intervals use 10,000 whole-exposure bootstrap resamples and a Bonferroni adjustment for six contrasts (99.1667% marginal percentile intervals). This is a bootstrap approximation, not a finite-sample coverage theorem. No model selection uses these outcomes.

Table 14. Frozen continuous-density and pose validation: relative half-width ranges over both targets and all four geometry replicates per stack. All 96 settings and all in-class signals are retained. The density, noise and pose limits are supplied assumptions.

EMPIAR	Scale 0.03		Scale 0.07	
	0.1°	0.5°	0.1°	0.5°
10028	0.76–0.77	0.98–1.00	0.11–0.14	0.31–0.38
10049	0.77–0.78	0.93–0.95	0.14–0.17	0.38–0.46
10076	0.76–0.77	0.89–0.89	0.12–0.14	0.27–0.31




Figure 7. All 48 frozen integrated-moment settings on disjoint particle subsets with new simulation seeds. The original spatial-maximum pose correction falls back to no data throughout; integrated moments retain informative widths in 32 settings. Conservative coverage and poor cell-generator sign power remain limitations. “L2” denotes densities at the declared ball boundary, chosen in the continuous adjoint-residual direction to maximize bias at a fixed pose.




Figure 8. Post-hoc integrated-moment refinement for both broad targets and every tested pose radius. Curves show uncapped widths to expose conservatism; the implemented decision uses the no-data estimator above the dashed line. The solid curves use integrated Fourier-column moments and actual estimator weights; dotted curves use the older spatial-maximum cubic remainder. No weights are refitted.




Figure 9. All 18 continuous nonlinear stress settings, retaining the strongest of six feasible searches per setting. The left panel compares numerical feasible bias with the integrated-moment upper bound; the right compares it with the earlier coherent-pose control. Surrogate optimization proposes poses only: final candidates eliminate the full continuous density ball using analytic sinc integrals. The largest independent direct-field relative discrepancy was 5.35×10^{-6} . Remaining upper/lower gaps and local-search limitations preclude a sharpness claim.




Figure 10. Descriptive background dependence in every additional-exposure image, after Fourier cropping of the downloaded particles to the working grid and before reconstruction windowing. We center and scale each image using pixels outside radius 0.43 of the box, then average products of background pixel pairs separated by each lag. Error bars are marginal 95% whole-exposure bootstrap intervals. Dotted lines show matched-statistic white-Gaussian controls. Background contamination and preprocessing prevent interpreting these values as pure-noise correlations or a confidence calibration.




Figure 11. Left: optimized continuous fixed-pose widths for the full cube class. Right: cell-projected audit coverage at a constructed continuous bias boundary, for the broad central target and cell-fitted estimator. Reference-map coverage is one in these cases. The two panels answer distinct questions: useful width and omitted continuous directions.




Figure 12. Continuous-density and nonlinear-pose development audit at target width 0.07 of the field. Solid curves are uniform conditional upper bounds for the same fixed-pose-optimized weights. Dashed curves use exact density-ball elimination at one allowed coherent rotation, providing feasible lower bounds on the required width, not global pose optima. Zero is the separate exact-pose control; the nonzero budgets also allow shifts of 0.01/24 field units. The visible gap exposes conservatism.




Figure 13. Deliberate assumption stresses for the finite-voxel central target. Curves share a signal within each geometry/control: the fixed-pose boundary for noise tests, a fixed-pose-estimator-directed support or mixture perturbation, or a saved nonlinear pose adversary. Dashed curves are fixed-pose certificates; solid curves include the shared-density pose audit. These are constructed diagnostic conditions, not experimental failure rates. All other controls remain in the result records.




Figure 14. Exact-dictionary development tests with known bounds and an independent fixed pilot. Minimum coverage is over the finite tested set, including adversarial signals; it is not an empirical proof of the worst case. Widths count each independent design once. Gaussian posterior methods separately pass their correctly matched prior-predictive coverage check at 0.95. Their fixed-parameter boundary failures do not contradict Bayesian prior-average calibration.




Figure 15. Central-average coverage and widths across all four ridge penalties. Each geometry uses its own fixed independent deposited-map generator. Error bars are binomial 95% intervals over independent measurement-noise draws, conditional on the reused group-bootstrap plan. Weak-regularization bootstrap can be narrow and near nominal; stronger regularization exposes bias. These tests do not establish empirical coverage for unknown experimental density.




Figure 16. The same complete regularization comparison for the axial contrast. Both the estimand and the unknown simulation signal are fixed before noise replication. The full results also retain the smaller 33-mode dictionary, basic bootstrap and Gaussian covariance controls.




Figure 17. Development tuning error at all evaluated neural checkpoints, with the classical reconstruction controls. Checkpoint selection minimizes this tuning error; the test-group error is not used to select the checkpoint. Different training objectives and supplied consensus poses limit this comparison.